

GLMM, Concepts, & R

Bill Last Updated:

18 January, 2020

Contents

Preface: Motivation	7
1 Basics	9
1.1 Logit	9
1.2 Probit	11
2 MLE	15
2.1 Basic idea of MLE	15
2.2 Coin flip example, probit, and logit	16
2.3 Further on logit	17
2.4 References	19
3 Linear Mixed Models	21
3.1 LM and GLM	21
3.2 LMM	30
3.3 Calculate mean	30
3.4 Test the treatment effect	32
3.5 Another example	33
3.6 Full LMM model	34
3.7 Serial correlations in time and space	36
4 Basic Stat Concepts	41
4.1 Score	41
4.2 Gradient and Jacobian	42
4.3 Hessian and Fisher Information	42
4.4 Canonical link function	43
4.5 Ordinary Least Squares (OLS)	44
4.6 Taylor series	46
4.7 Fisher scoring	46
4.8 References	47
5 Basic R	49
5.1 apply, lapply, sapply	49
5.2 C	50

6	Computing Techniques	53
6.1	Monte carlo approximation	53
6.2	Importance sampling	54
6.3	Newton Raphson algorithm	56
6.4	Metropolis Hastings	62
6.5	EM	64
6.6	References	64
7	Generalized Linear Mixed Models	65
7.1	Basics of GLMM	65
7.2	Some References	66
8	Twitter Example	67
8.1	Model	67
8.2	Simulating Data of Senators on Twitter	69
8.3	Simulating Data of Conservative Users on Twitter and Model Testing	70
8.4	Simulating Data of Liberal Users on Twitter and Model Testing	72
9	Practice: Learning on the Battle Field	75
9.1	R code	75
9.2	References	80
10	Project Draft	81
10.1	Background	83
10.2	Important Examples with R code	86
10.3	References	86
11	Bayesian - 1	87
11.1	Frequentist perspective	87
11.2	Bayesian perspective	88
11.3	Continuous parameters	88
11.4	Bernoulli/binomial likelihood with uniform prior	90
11.5	Conjugate priors	91
11.6	Poisson distribution	93
11.7	Exponential data	94
11.8	Normal likelihood	96
11.9	Non-informative priors	97
11.10	Jeffreys Priors	99
11.11	Linear regression	100
12	Bayesian - 2	103
12.1	Components of Bayesian models	103
12.2	Monte Carlo Estimation	104
12.3	Prior predictive distributions	106
13	Trying	107

<i>CONTENTS</i>	5
13.1 The Basic Idea	107
13.2 Model and R Code	107
13.3 glmmTMB package	112

Preface: Motivation

All the notes I have done here are the preparation for my stat master project, which will be about Generalized Linear Mixed Models. While I have tried my best, probably there are still some typos and errors. Please feel free to let me know in case you find one. Thank you!

Chapter 1

Basics

1.1 Logit


$$f(x) = \log\left(\frac{p(y=1)}{1-p(y=1)}\right)$$

The basic idea of logistic regression:

$$p(y=1) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n)}} = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}}$$

Thus, $e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}$ can be from $-\infty$ to $+\infty$, and $p(y=1)$ will be always within the range of $(0, 1)$.

```
f<-function(x){exp(x)/(1+exp(x))}  
data<-seq(-10,10,1)  
plot(data,f(data),type = "b")
```



We can also write the function into another format as follows:

$$\log \frac{p(y=1)}{1-p(y=1)} = \beta_0 + \beta_1 x_1 + \dots + \beta_n x_n$$

Thus, we know that the regression coefficients of β_i actually change the “log-odds” of the event. Of course, note that the magnitude of β_i is dependent upon the units of x_i .

The following is an example testing whether that home teams are more likely to win in NFL games. The results show that the odd of winning is the same for both home and away teams.

```
mydata = read.csv(url('https://raw.githubusercontent.com/nfl-football-ops/Big-Data-Bow'))
mydata$result_new<-ifelse(mydata$HomeScore>mydata$VisitorScore,1,0)
summary(mydata$result_new)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## 0.0000 0.0000 0.0000 0.4945 1.0000 1.0000
```

```
mylogit1 = glm(result_new~1, family=binomial, data=mydata)
summary(mylogit1)
```


```
##
## Call:
## glm(formula = result_new ~ 1, family = binomial, data = mydata)
##
## Deviance Residuals:
```

```
##      Min      1Q  Median      3Q      Max
## -1.168 -1.168 -1.168   1.187   1.187
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.02198    0.20967  -0.105   0.917
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 126.14  on 90  degrees of freedom
## Residual deviance: 126.14  on 90  degrees of freedom
## AIC: 128.14
##
## Number of Fisher Scoring iterations: 3
```

1.2 Probit

As noted above, logit $f(x) = \log\left(\frac{p(y=1)}{1-p(y=1)}\right)$ provides the resulting range of $(0,1)$. Another way to provide the same range is through the cdf of normal distribution. The following R code is used to illustrate this process.

```
data2<-seq(-5,5,1)
plot(data2,pnorm(data2),type = "b")
```



Thus, the cdf of normal distribution can be used to indicate the probability of $p(y = 1)$.

$$\Phi(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n) = p(y = 1)$$

Similar to logit model, we can also write the inverse function of the cdf to get the function that can be from $-\infty$ to $+\infty$.

$$\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n = \Phi^{-1}(p(y = 1))$$

Thus, for example, if $X\beta = -2$, based on $\Phi(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n) = p(y = 1)$ we can get that the $p(y = 1) = 0.023$.

In contrast, if $X\beta = 3$, the $p(y = 1) = 0.999$.

```
pnorm(-2)
```

```
## [1] 0.02275013
```

```
pnorm(3)
```

```
## [1] 0.9986501
```

Let's assume that there is a latent variable called Y^* such that

$$Y^* = X\beta + \epsilon, \epsilon \sim N(0, \sigma^2)$$

You could think of Y^* as a kind of “proxy” between $X\beta + \epsilon$ and the observed $Y(1 \text{ or } 0)$. Thus, we can get the following. Note that, it does not have to be zero, and can be any constant.

$$Y^* = \begin{cases} 0 & \text{if } y_i^* \leq 0 \\ 1 & \text{if } y_i^* > 0 \end{cases}$$

Thus,

$$y_i^* > 0 \Rightarrow \beta' X_i + \epsilon_i > 0 \Rightarrow \epsilon_i > -\beta' X_i$$

Thus, we can write it as follows. Note that $\frac{\epsilon_i}{\sigma} \sim N(0, 1)$

$$p(y = 1|x_i) = p(y_i^* > 0|x_i) = p(\epsilon_i > -\beta' X_i) = p\left(\frac{\epsilon_i}{\sigma} > \frac{-\beta' X_i}{\sigma}\right) = \Phi\left(\frac{\beta' X_i}{\sigma}\right)$$

We thus can get:

$$p(y = 0|x_i) = 1 - \Phi\left(\frac{\beta' X_i}{\sigma}\right)$$

For $p(y = 1|x_i) = \Phi(\frac{\beta' X_i}{\sigma})$, we can not really estimate both β and σ as they are in a ratio. We can assume $\sigma = 1$, then $\epsilon \sim N(0, 1)$. We know y_i and x_i since we observe them. Thus, we can write it as follows.

$$p(y = 1|x_i) = \Phi(\beta' X_i)$$

Chapter 2

MLE

2.1 Basic idea of MLE

Suppose that we flip a coin, $y_i = 0$ for tails and $y_i = 1$ for heads. If we get p heads from n trials, we can get the proportion of heads is p/n , which is the sample mean. If we do not do any further calculation, this is our best guess.

Suppose that the true probability is ρ , then we can get:

$$\mathbf{L}(y_i) = \begin{cases} \rho & y_i = 1 \\ 1 - \rho & y_i = 0 \end{cases}$$

Thus, we can also write it as follows.

$$\mathbf{L}(y_i) = \rho^{y_i} (1 - \rho)^{1-y_i}$$

Thus, we can get:

$$\prod \mathbf{L}(y_i|\rho) = \rho^{\sum y_i} (1 - \rho)^{\sum (1-y_i)}$$

Further, we can get a log-transformed format.

$$\log(\prod \mathbf{L}(y_i|\rho)) = \sum y_i \log \rho + \sum (1 - y_i) \log(1 - \rho)$$

To maximize the log-function above, we can calculate the derivative with respect to ρ .

$$\frac{\partial \log(\prod \mathbf{L}(y_i|\rho))}{\partial \rho} = \sum y_i \frac{1}{\rho} - \sum (1 - y_i) \frac{1}{1 - \rho}$$

Set the derivative to zero and solve for ρ , we can get

$$\begin{aligned}
& \sum y_i \frac{1}{\rho} - \sum (1 - y_i) \frac{1}{1 - \rho} = 0 \\
& \Rightarrow (1 - \rho) \sum y_i - \rho \sum (1 - y_i) = 0 \\
& \Rightarrow \sum y_i - \rho \sum y_i - n\rho + \rho \sum y_i = 0 \\
& \Rightarrow \sum y_i - n\rho = 0 \\
& \Rightarrow \rho = \frac{\sum y_i}{n} = \frac{p}{n}
\end{aligned}$$

Thus, we can see that the ρ maximizing the likelihood function is equal to the sample mean.

2.2 Coin flip example, probit, and logit

In the example above, we are not really trying to estimate a lot of regression coefficients. What we are doing actually is to calculate the sample mean, or intercept in the regression sense. What does it mean? Let's use some data to explain it.

Suppose that we flip a coin 20 times and observe 8 heads. We can use the R's `glm` function to estimate the ρ . If the result is consistent with what we did above, we should observe that the *cdf* of the estimate of β_0 (i.e., intercept) should be equal to $8/20 = 0.4$.

```
coins<-c(rep(1,times=8),rep(0,times=12))
table(coins)
```

```
## coins
##  0  1
## 12  8
```

```
coins<-as.data.frame(coins)
```

2.2.1 Probit

```
probitresults <- glm(coins ~ 1, family = binomial(link = "probit"), data = coins)
probitresults
```

```
##
## Call:  glm(formula = coins ~ 1, family = binomial(link = "probit"),
##       data = coins)
##
## Coefficients:
## (Intercept)
##      -0.2533
```



```
##
## Degrees of Freedom: 19 Total (i.e. Null); 19 Residual
## Null Deviance: 26.92
## Residual Deviance: 26.92 AIC: 28.92
```

```
pnorm(probitresults$coefficients)
```

```
## (Intercept)
## 0.4
```

As we can see the intercept is -0.2533 , and thus $\Phi(-0.2533471) = 0.4$

2.2.2 Logit

We can also use logit link to calculate the intercept as well. Recall that

$$p(y = 1) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n)}} = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}}$$

Thus,

$$p(y = 1) = \frac{e^{\beta_0}}{1 + e^{\beta_0}}$$

```
logitresults <- glm(coins ~ 1, family = binomial(link = "logit"), data = coins)
logitresults$coefficients
```

```
## (Intercept)
## -0.4054651
```

```
exp(logitresults$coefficients)/(1+exp(logitresults$coefficients))
```

```
## (Intercept)
## 0.4
```

Note that, the default link for the binomial in the glm function is logit.

2.3 Further on logit

The probability of $y = 1$ is as follows:

$$p = p(y = 1) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n)}} = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}}$$

Thus, the likelihood function is as follows:

$$L = \prod p^{y_i} (1-p)^{1-y_i} = \prod \left(\frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n)}} \right)^{y_i} \left(\frac{1}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}} \right)^{1-y_i}$$

$$= \prod (1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n)})^{-y_i} (1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n})^{-(1-y_i)}$$

Thus, the log-likelihood is as follows:

$$\log L = \sum (-y_i \cdot \log(1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n)}) - (1 - y_i) \cdot \log(1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}))$$

Typically, optimisers minimize a function, so we use negative log-likelihood as minimising that is equivalent to maximising the log-likelihood or the likelihood itself.

#Source of R code: <https://www.r-bloggers.com/logistic-regression/>

```
mle.logreg = function(fmla, data)
{
  # Define the negative log likelihood function
  logl <- function(theta,x,y){
    y <- y
    x <- as.matrix(x)
    beta <- theta[1:ncol(x)]

    # Use the log-likelihood of the Bernoulli distribution, where p is
    # defined as the logistic transformation of a linear combination
    # of predictors, according to logit(p)=(x%*%beta)
    loglik <- sum(-y*log(1 + exp(-(x%*%beta))) - (1-y)*log(1 + exp(x%*%beta)))
    return(-loglik)
  }

  # Prepare the data
  outcome = rownames(attr(terms(fmla),"factors"))[1]
  dfrTmp = model.frame(data)
  x = as.matrix(model.matrix(fmla, data=dfrTmp))
  y = as.numeric(as.matrix(data[,match(outcome,colnames(data))]))

  # Define initial values for the parameters
  theta.start = rep(0,(dim(x)[2]))
  names(theta.start) = colnames(x)

  # Calculate the maximum likelihood
  mle = optim(theta.start,logl,x=x,y=y, method = 'BFGS', hessian=T)
  out = list(beta=mle$par,vcov=solve(mle$hessian),ll=2*mle$value)
}

mydata = read.csv(url('https://stats.idre.ucla.edu/stat/data/binary.csv'))
mylogit1 = glm(admit~gre+gpa+as.factor(rank), family=binomial, data=mydata)
```

```
mydata$rank = factor(mydata$rank) #Treat rank as a categorical variable
fmla = as.formula("admit~gre+gpa+rank") #Create model formula
mylogit2 = mle.logreg(fmla, mydata) #Estimate coefficients

print(cbind(coef(mylogit1), mylogit2$beta))
```

```
##                [,1]      [,2]
## (Intercept)    -3.989979073 -3.772676422
## gre             0.002264426  0.001375522
## gpa             0.804037549  0.898201239
## as.factor(rank)2 -0.675442928 -0.675543009
## as.factor(rank)3 -1.340203916 -1.356554831
## as.factor(rank)4 -1.551463677 -1.563396035
```

2.4 References

http://www.columbia.edu/~so33/SusDev/Lecture_9.pdf

Chapter 3

Linear Mixed Models

3.1 LM and GLM

Before moving to LMM, I would like to review LM and GLM first.

3.1.1 LM

$$Y|X \sim N(\mu(X), \sigma^2 I)$$

$$E(Y|X) = \mu(X) = X^T \beta$$

where,

$\mu(X)$: *random component*

$X^T \beta$: *covariates*

3.1.2 GLM-Definition

Ref: https://ocw.mit.edu/courses/mathematics/18-650-statistics-for-applications-fall-2016/lecture-slides/MIT18_650F16_GLM.pdf

$$Y \sim \text{exponential family}$$

Link function

$$g(\mu(X)) = X^T \beta$$

3.1.3 GLM-log link example

$$\mu_i = \gamma e^{\delta t_i}$$

Link function is log link, and it becomes:

$$\log(\mu_i) = \log(\gamma) + \log(\delta t_i) = \beta_0 + \beta_1 t_i$$

(This is somehow similar to Poisson distribution.)

3.1.4 GLM-Reciprocal link:

$$\mu_i = \frac{\alpha x_i}{h + x_i}$$

Reciprocal link:

$$g(\mu_i) = \frac{1}{\mu_i} = \frac{1}{\alpha} + \frac{h}{\alpha} \frac{1}{x_i} = \beta_0 + \beta_1 \frac{1}{x_i}$$

3.1.5 GLM-exponential family:

In a more general sense, for exponential family:

$$\begin{aligned} P_\theta(X) &= P(X, \theta) = e^{\sum \eta_i(\theta) T_i(X)} C(\theta) h(x) \\ &= e^{\sum \eta_i(\theta) T_i(X)} e^{-\log(\frac{1}{c(\theta)})} h(x) \\ &= e^{\sum \eta_i(\theta) T_i(X) - \log(\frac{1}{c(\theta)})} h(x) \\ &= e^{\sum \eta_i(\theta) T_i(X) - B(\theta)} h(x) \end{aligned}$$

Normal distribution

For normal distributions, it belongs to exponential family.

$$\begin{aligned} P_\theta(X) &= \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}} \\ &= e^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}} e^{\log(\frac{1}{\sigma\sqrt{2\pi}})} \\ &= e^{-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2} - \log(\sigma\sqrt{2\pi})} \\ &= e^{-\frac{1}{2\sigma^2} x^2 - \frac{1}{2\sigma^2} \mu^2 + \frac{x\mu}{\sigma^2} - \log(\sqrt{2\pi}\sigma)} \\ &= e^{-\frac{1}{2\sigma^2} x^2 + \frac{x\mu}{\sigma^2} - (\frac{1}{2\sigma^2} \mu^2 + \log(\sqrt{2\pi}\sigma))} \end{aligned}$$

Where,

$$\eta_1 = -\frac{1}{2\sigma^2} \text{ and } T_1(x) = x^2$$

$$\eta_2 = -\frac{\mu}{\sigma^2} \text{ and } T_2(x) = x$$

$$B(\theta) = \frac{1}{2\sigma^2}\mu^2 + \log(\sqrt{2\pi}\sigma)$$

$$h(x) = 1$$

In the case above, $\theta = (\mu, \sigma^2)$. If σ^2 is known, $\theta = \mu$. In this case, we can rewrite the normal pdf as follows.

$$\begin{aligned} P_\theta(X) &= e^{-\frac{1}{2\sigma^2}x^2 - \frac{1}{2\sigma^2}\mu^2 + \frac{x\mu}{\sigma^2} - \log(\sqrt{2\pi}\sigma)} \\ &= e^{\frac{x\mu}{\sigma^2} - \frac{1}{2\sigma^2}\mu^2} e^{-\frac{1}{2\sigma^2}x^2 - \log(\sqrt{2\pi}\sigma)} \end{aligned}$$

Where,

$$\eta_1 = -\frac{\mu}{\sigma^2} \text{ and } T_1(x) = x$$

$$B(\theta) = \frac{1}{2\sigma^2}\mu^2$$

$$\begin{aligned} h(x) &= e^{-\frac{1}{2\sigma^2}x^2 - \log(\sqrt{2\pi}\sigma)} \\ &= \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\frac{x^2}{\sigma^2}} \end{aligned}$$

Thus, we can see that $h(x)$ is a normal pdf $\sim N(0, \sigma^2)$.

Bernoulli

Another example, x is discrete. For example, Bernoulli:

$$\begin{aligned} &= p^x(1-p)^{1-x} \\ &= e^{\log(p^x(1-p)^{1-x})} \\ &= e^{x\log(p) + (1-x)\log(1-p)} \\ &= e^{x\log(p) - x\log(1-p) + \log(1-p)} \\ &= e^{x\log(\frac{p}{1-p}) + \log(1-p)} \end{aligned}$$

Where,

$$\eta_1 = \log\left(\frac{p}{1-p}\right) \text{ and } T_1(x) = x$$

$$B(\theta) = \log\left(\frac{1}{1-p}\right)$$

$$h(x) = 1$$

3.1.6 Canonical exponential family

Canonical exponential family:

$$f_\theta(x) = e^{\frac{x\theta - b(\theta)}{\phi} + c(x, \phi)}$$

where, $b(\cdot)$ and $c(\cdot, \cdot)$ are known.

Normal distribution

Again, use the normal pdf:

$$\begin{aligned}
 P_{\theta}(X) &= \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}} \\
 &= e^{-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2}} e^{\log(\frac{1}{\sigma\sqrt{2\pi}})} \\
 &= e^{-\frac{1}{2}\frac{(x-\mu)^2}{\sigma^2} - \log(\sigma\sqrt{2\pi})} \\
 &= e^{-\frac{1}{2\sigma^2}x^2 - \frac{1}{2\sigma^2}\mu^2 + \frac{x\mu}{\sigma^2} - \log(\sqrt{2\pi}\sigma)} \\
 &= e^{\frac{x\mu}{\sigma^2} - \frac{\mu^2}{2\sigma^2} + (-\frac{1}{2\sigma^2}x^2 - \log(\sqrt{2\pi}\sigma))} \\
 &= e^{\frac{x\mu - \frac{1}{2}\mu^2}{\sigma^2} + (-\frac{1}{2\sigma^2}x^2 - \log(\sqrt{2\pi}\sigma))}
 \end{aligned}$$

Where (we assume σ^2 is known.),

$$\theta = \mu$$

$$\phi = \sigma^2$$

$$b(\theta) = \frac{1}{2}\theta^2$$

$$\begin{aligned}
 c(x, \phi) &= -\frac{1}{2\sigma^2}x^2 - \log(\sqrt{2\pi}\sigma) \\
 &= -\frac{1}{2\sigma^2}x^2 - \frac{1}{2}\log(2\pi\sigma^2) \\
 &= -\frac{1}{2}\left(\frac{x^2}{\sigma^2} + \log(2\pi\sigma^2)\right) \\
 &= -\frac{1}{2}\left(\frac{x^2}{\phi} + \log(2\pi\phi)\right)
 \end{aligned}$$

3.1.7 Canonical exponential family - Expected value and variance**First derivative**

Canonical exponential family:

$$f_{\theta}(x) = e^{\frac{x\theta - b(\theta)}{\phi} + c(x, \phi)}$$

log likelihood (only one observation)

$$\log f_{\theta}(x)$$

$$\begin{aligned}
E\left[\frac{\partial(\log f_\theta(X))}{\partial\theta}\right] &= E\left[\frac{\frac{\partial f_\theta(X)}{\partial\theta}}{f_\theta(X)}\right] \\
&= \int \frac{\frac{\partial f_\theta(X)}{\partial\theta}}{f_\theta(X)} f_\theta(X) dx \\
&= \int \frac{\partial f_\theta(X)}{\partial\theta} dx \\
&= \frac{\partial}{\partial\theta} \int f_\theta(X) dx \\
&= \frac{\partial 1}{\partial\theta} \\
&= 0
\end{aligned}$$

Second derivative

Second derivative

$$\begin{aligned}
E\left[\frac{\partial^2(\log f_\theta(X))}{\partial\theta^2}\right] &= E\left[\frac{\partial}{\partial\theta}\left(\frac{\frac{\partial f_\theta(X)}{\partial\theta}}{f_\theta(X)}\right)\right] \\
&= E\left[\frac{\frac{\partial^2 f_\theta(X)}{\partial\theta^2} f_\theta(X) - \left(\frac{\partial f_\theta(X)}{\partial\theta}\right)^2}{f_\theta^2(X)}\right] \\
&= \int \frac{\frac{\partial^2 f_\theta(X)}{\partial\theta^2} f_\theta(X) - \left(\frac{\partial f_\theta(X)}{\partial\theta}\right)^2}{f_\theta^2(X)} dx \\
&= \int \left(\frac{\partial^2 f_\theta(X)}{\partial\theta^2} - \frac{\left(\frac{\partial f_\theta(X)}{\partial\theta}\right)^2}{f_\theta(X)}\right) dx \\
&= \int \frac{\partial^2 f_\theta(X)}{\partial\theta^2} dx - \int \frac{\left(\frac{\partial f_\theta(X)}{\partial\theta}\right)^2}{f_\theta(X)} dx \\
&= \frac{\partial^2}{\partial\theta^2} \int f_\theta(X) dx - \int \frac{\left(\frac{\partial f_\theta(X)}{\partial\theta}\right)^2}{f_\theta(X)} dx \\
&= 0 - \int \frac{\left(\frac{\partial f_\theta(X)}{\partial\theta}\right)^2}{f_\theta(X)} dx \\
&= 0 - \int \frac{\left(\frac{\partial f_\theta(X)}{\partial\theta}\right)^2}{(f_\theta(X))^2} f_\theta(x) dx \\
&= -E\left[\left(\frac{\frac{\partial f_\theta(X)}{\partial\theta}}{f_\theta(X)}\right)^2\right] \\
&= -E\left[\left(\frac{\partial(\log f_\theta(X))}{\partial\theta}\right)^2\right]
\end{aligned}$$

Based on the first derivative, we can get:

$$\log(f_\theta(X)) = \frac{X\theta - b(\theta)}{\phi} + c(X, \phi)$$

$$E\left[\frac{\partial(\log(f_\theta(X)))}{\partial\theta}\right] = E\left[\frac{X - b'(\theta)}{\phi}\right] = \frac{E(X) - b'(\theta)}{\phi} = 0$$

Thus, we can get,

$$E(X) = b'(\theta)$$

For second derivative, from the calculation above, we know that,

$$\begin{aligned} E\left[\frac{\partial^2(\log f_\theta(X))}{\partial\theta^2}\right] &= -E\left[\left(\frac{\partial(\log f_\theta(X))}{\partial\theta}\right)^2\right] \\ &= -E\left[\left(\frac{X - b'(\theta)}{\phi}\right)^2\right] \\ &= -E\left[\left(\frac{X - E(X)}{\phi}\right)^2\right] \\ &= -\frac{\text{Var}(X)}{\phi^2} \end{aligned}$$

At the same time,

$$\begin{aligned} E\left[\frac{\partial^2(\log f_\theta(X))}{\partial\theta^2}\right] &= E\left[\frac{\partial\left(\frac{X - b'(\theta)}{\phi}\right)}{\partial\theta}\right] \\ &= E\left[-\frac{b''(\theta)}{\phi}\right] \\ &= -\frac{b''(\theta)}{\phi} \end{aligned}$$

Thus,

$$\text{Var}(X) = b''(\theta)\phi$$

3.1.8 Expected value and variance - Poisson Example

Example of poisson distribution

$$P(\lambda) = \frac{\lambda^k e^{-\lambda}}{k!}$$

If we put k as y , and λ as μ , we can get:

$$P(\mu) = \frac{\mu^y e^{-\mu}}{y!}$$

Compare to,

$$f_{\theta}(y) = e^{\frac{y\theta - b(\theta)}{\phi} + c(y, \phi)}$$

We can write it as the exponential format:

$$\begin{aligned} P(\mu) &= \frac{\mu^y e^{-\mu}}{y!} \\ &= e^{\log(\mu^y) + \log(e^{-\mu}) - \log(y!)} \\ &= e^{y \log(\mu) - \mu - \log(y!)} \end{aligned}$$

We thus know that $\theta = \log(\mu)$. We can continue to write the equation above as follows.

$$= e^{y\theta - e^{\theta} - \log(y!)}$$

Thus, we can get:

$$\begin{aligned} E(X) &= \frac{\partial(e^{\theta})}{\partial\theta} = e^{\theta} = \mu \\ Var(X) &= \frac{\partial''(e^{\theta})}{\partial\theta^2} \phi = \frac{\partial''(e^{\theta})}{\partial\theta^2} = \mu \end{aligned}$$

3.1.9 Canonical link

A link function can link $X^T\beta$ to the mean μ .

That is,

$$g(\mu) = X^T\beta \rightarrow \mu = g^{-1}(X^T\beta)$$

We know that

$$\mu = b'(\theta)$$

Thus,

$$b'(\theta) = g^{-1}(X^T\beta)$$

Thus,

$$g = b'^{-1}(\theta)$$

3.1.10 Canonical link - Bernoulli

PMF of Bernoulli:

$$\begin{aligned} &= p^y(1-p)^{1-y} \\ &= e^{\log(p^y(1-p)^{1-y})} \\ &= e^{y\log(p)+(1-y)\log(1-p)} \\ &= e^{y\log(p)-y\log(1-p)+\log(1-p)} \\ &= e^{y\log(\frac{p}{1-p})+\log(1-p)} \end{aligned}$$

Compared to the following:

$$f_{\theta}(y) = e^{\frac{y\theta - b(\theta)}{\phi} + c(y, \phi)}$$

We need to change the format of Bernoulli:

$$\theta = \log \frac{p}{1-p}$$

Thus,

$$e^{\theta} = \frac{p}{1-p} \rightarrow p = \frac{e^{\theta}}{1+e^{\theta}}$$

After that, we can continue the Bernoulli:

$$\begin{aligned} &= e^{y\theta + \log(1 - \frac{e^{\theta}}{1+e^{\theta}})} \\ &= e^{y\theta + \log(\frac{1}{1+e^{\theta}})} \\ &= e^{y\theta - \log(1+e^{\theta})} \end{aligned}$$

Where,

$$b(\theta) = \log(1 + e^{\theta})$$

We can then try to calculate the derivative:

$$b'(\theta) = \frac{\partial(\log(1 + e^{\theta}))}{\partial \theta} = \frac{e^{\theta}}{1 + e^{\theta}}$$

We know that

$$b'(\theta) = \mu$$

Thus, we can get

$$\mu = \frac{e^\theta}{1 + e^\theta}$$

We can then calculate the inverse function:

$$\theta = \log\left(\frac{\mu}{1 - \mu}\right)$$

Thus,

$$g(\mu) = \log\left(\frac{\mu}{1 - \mu}\right) = X^T \beta$$

3.1.11 NR - Bernoulli

We know that the PMF for Bernoulli:

$$\begin{aligned} &= p^y (1 - p)^{1-y} \\ &= e^{y\theta - \log(1 + e^\theta)} \\ &= e^{yx^T \beta - \log(1 + e^{x^T \beta})} \end{aligned}$$

Thus,

$$\ell(\beta|Y, X) = \sum_{i=1}^n (Y_i X_i^T \beta - \log(1 + e^{X_i^T \beta}))$$

Thus, the gradient is:

$$\nabla_\ell(\beta) = \sum_{i=1}^n \left(Y_i X_i - \frac{e^{X_i^T \beta}}{1 + e^{X_i^T \beta}} \right)$$

The Hessian is:

$$H_\ell(\beta) = - \sum_{i=1}^n \frac{e^{X_i^T \beta}}{(1 + e^{X_i^T \beta})^2} X_i X_i^T$$

Thus,

$$\beta^{k+1} = \beta^k - (H_\ell(\beta^k))^{-1} \nabla_\ell(\beta^k)$$

3.1.12 Iteratively Re-weighted Least Squares

3.2 LMM

The following is a shortened version of Jonathan Rosenblatt's LMM tutorial. <http://www.john-ros.com/Rcourse/lme.html>.

In addition, another reference is from Douglas Bates's R package document. https://cran.r-project.org/web/packages/lme4/vignettes/lmer.pdf?fbclid=IwAR1nmmRP9A0BrhKdgBibNjM5acR_spTpXV8QlQGdmTWyQz3ZtV3LYn6kCbQ

Assume that y is a function of x and u , where x is the fixed effect and u is the random effect. Thus, we can get,

$$y|x, u = x'\beta + z'u + \epsilon$$

For random effect, one example can be that you want to test the treatment effect, and sample 8 observations from 4 groups. You measure before and after the treatment. In this case, x represents the treatment effect, whereas z represents the group effect (i.e., random effect). Note that, in this case, it reminds the paired t-test. Remember in SPSS, why do we do paired t-test? Typically, it is the case when we measure a subject (or, participant) twice. In this case, we can consider each participant as an unit of random effect (rather than as group in the last example.)

3.3 Calculate mean

The following code generates 4 numbers ($N(0, 10)$) for 4 groups. Then, replicate it within each group. That is, in the end, there are 8 observations.

Note that, in the following code, there are no "independent variables". Both the linear model and mixed model are actually just trying to calculate the mean. Note that `lmer(y~1+1|groups)` and `lmer(y~1|groups)` will generate the same results.

```
set.seed(123)
n.groups <- 4 # number of groups
n.repeats <- 2 # samples per group
#Generating index for observations belong to the same group
groups <- as.factor(rep(1:n.groups, each=n.repeats))
n <- length(groups)
#Generating 4 random numbers, assuming normal distribution
z0 <- rnorm(n.groups, 0, 10)
z <- z0[as.numeric(groups)] # generate and inspect random group effects
z

## [1] -5.6047565 -5.6047565 -2.3017749 -2.3017749 15.5870831 15.5870831
```

```
## [7] 0.7050839 0.7050839

epsilon <- rnorm(n,0,1) # generate measurement error
beta0 <- 2 # this is the actual parameter of interest! The global mean.
y <- beta0 + z + epsilon # sample from an LMM

# fit a linear model assuming independence
# i.e., assume that there is no "group things".
lm.5 <- lm(y~1)

# fit a mixed-model that deals with the group dependence
#install.packages("lme4")
library(lme4)
lme.5.a <- lmer(y~1+1|groups)
lme.5.b <- lmer(y~1|groups)
lm.5

##
## Call:
## lm(formula = y ~ 1)
##
## Coefficients:
## (Intercept)
##      4.283

lme.5.a

## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ 1 + 1 | groups
## REML criterion at convergence: 36.1666
## Random effects:
## Groups      Name          Std.Dev.
## groups      (Intercept) 8.8521
## Residual                0.8873
## Number of obs: 8, groups: groups, 4
## Fixed Effects:
## (Intercept)
##      4.283

lme.5.b

## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ 1 | groups
## REML criterion at convergence: 36.1666
## Random effects:
## Groups      Name          Std.Dev.
## groups      (Intercept) 8.8521
## Residual                0.8873
```

```
## Number of obs: 8, groups: groups, 4
## Fixed Effects:
## (Intercept)
##          4.283
```

3.4 Test the treatment effect

As we can see that, LLM and paired t-test generate the same t-value.

```
times<-rep(c(1,2),4) # first time and second time
times

## [1] 1 2 1 2 1 2 1 2

data_combined<-cbind(y,groups,times)
data_combined

##           y groups times
## [1,] -3.4754687      1      1
## [2,] -1.8896915      1      2
## [3,]  0.1591413      2      1
## [4,] -1.5668361      2      2
## [5,] 16.9002303      3      1
## [6,] 17.1414212      3      2
## [7,]  3.9291657      4      1
## [8,]  3.0648977      4      2

lme_diff_times<- lmer(y~times+(1|groups))

t_results<-t.test(y~times, paired=TRUE)

lme_diff_times

## Linear mixed model fit by REML ['lmerMod']
## Formula: y ~ times + (1 | groups)
## REML criterion at convergence: 35.0539
## Random effects:
## Groups      Name      Std.Dev.
## groups      (Intercept) 8.845
## Residual                    1.013
## Number of obs: 8, groups: groups, 4
## Fixed Effects:
## (Intercept)      times
##      4.5691      -0.1908

print("The following results are from paired t-test")
```



```
## [1] "The following results are from paired t-test"
```

```
t_results$statistic
```

```
##          t
```

```
## 0.2664793
```

3.5 Another example

```
data(Dyestuff, package='lme4')
attach(Dyestuff)
```

```
## The following objects are masked from Dyestuff (pos = 6):
```

```
##
```

```
##      Batch, Yield
```

```
Dyestuff
```

```
##      Batch Yield
```

```
## 1      A  1545
```

```
## 2      A  1440
```

```
## 3      A  1440
```

```
## 4      A  1520
```

```
## 5      A  1580
```

```
## 6      B  1540
```

```
## 7      B  1555
```

```
## 8      B  1490
```

```
## 9      B  1560
```

```
## 10     B  1495
```

```
## 11     C  1595
```

```
## 12     C  1550
```

```
## 13     C  1605
```

```
## 14     C  1510
```

```
## 15     C  1560
```

```
## 16     D  1445
```

```
## 17     D  1440
```

```
## 18     D  1595
```

```
## 19     D  1465
```

```
## 20     D  1545
```

```
## 21     E  1595
```

```
## 22     E  1630
```

```
## 23     E  1515
```

```
## 24     E  1635
```

```
## 25     E  1625
```

```
## 26     F  1520
```

```
## 27     F  1455
```

```
## 28      F 1450
## 29      F 1480
## 30      F 1445

lme_batch<- lmer( Yield ~ 1 + (1|Batch) , Dyestuff )
summary(lme_batch)

## Linear mixed model fit by REML ['lmerMod']
## Formula: Yield ~ 1 + (1 | Batch)
##      Data: Dyestuff
##
## REML criterion at convergence: 319.7
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.4117 -0.7634  0.1418  0.7792  1.8296
##
## Random effects:
##   Groups   Name      Variance Std.Dev.
##   Batch    (Intercept) 1764     42.00
##   Residual                2451     49.51
## Number of obs: 30, groups:  Batch, 6
##
## Fixed effects:
##              Estimate Std. Error t value
## (Intercept)  1527.50     19.38    78.8
```

3.6 Full LMM model

In the following, I used the data from the package of lme4. For Days + (1 | Subject), it only has random intercept; in contrast, Days + (Days| Subject) has both random intercept and random slope for Days. Note that, random effects do not generate specific slopes for each level of Days, but rather just a variance of all the slopes.

Therefore, we can see that “Days + (Days| Subject)” and “Days + (1+Days| Subject)” generate the same results. For more discussion, you can refer to the following link: <https://www.jaredknowles.com/journal/2013/11/25/getting-started-with-mixed-effect-models-in-r>

```
data(sleepstudy, package='lme4')
attach(sleepstudy)
```

```
## The following objects are masked from sleepstudy (pos = 6):
##
##      Days, Reaction, Subject
```

```

fm1 <- lmer(Reaction ~ Days + (1 | Subject), sleepstudy)
summary(fm1)

## Linear mixed model fit by REML ['lmerMod']
## Formula: Reaction ~ Days + (1 | Subject)
## Data: sleepstudy
##
## REML criterion at convergence: 1786.5
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.2257 -0.5529  0.0109  0.5188  4.2506
##
## Random effects:
## Groups Name Variance Std.Dev.
## Subject (Intercept) 1378.2  37.12
## Residual 960.5 30.99
## Number of obs: 180, groups: Subject, 18
##
## Fixed effects:
## Estimate Std. Error t value
## (Intercept) 251.4051 9.7467 25.79
## Days 10.4673 0.8042 13.02
##
## Correlation of Fixed Effects:
## (Intr)
## Days -0.371

fm2<-lmer ( Reaction ~ Days + ( Days| Subject ) , data= sleepstudy )
summary(fm2)

## Linear mixed model fit by REML ['lmerMod']
## Formula: Reaction ~ Days + (Days | Subject)
## Data: sleepstudy
##
## REML criterion at convergence: 1743.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.9536 -0.4634  0.0231  0.4633  5.1793
##
## Random effects:
## Groups Name Variance Std.Dev. Corr
## Subject (Intercept) 611.90 24.737
## Days 35.08 5.923 0.07
## Residual 654.94 25.592

```

```
## Number of obs: 180, groups: Subject, 18
##
## Fixed effects:
##           Estimate Std. Error t value
## (Intercept) 251.405      6.824  36.843
## Days        10.467      1.546   6.771
##
## Correlation of Fixed Effects:
##      (Intr)
## Days -0.138

fm3<-lmer ( Reaction ~ Days + (1+Days| Subject ) , data= sleepstudy )
summary(fm3)

## Linear mixed model fit by REML ['lmerMod']
## Formula: Reaction ~ Days + (1 + Days | Subject)
##   Data: sleepstudy
##
## REML criterion at convergence: 1743.6
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -3.9536 -0.4634  0.0231  0.4633  5.1793
##
## Random effects:
##   Groups   Name      Variance Std.Dev. Corr
##   Subject (Intercept) 611.90   24.737
##           Days        35.08    5.923   0.07
##   Residual          654.94   25.592
## Number of obs: 180, groups: Subject, 18
##
## Fixed effects:
##           Estimate Std. Error t value
## (Intercept) 251.405      6.824  36.843
## Days        10.467      1.546   6.771
##
## Correlation of Fixed Effects:
##      (Intr)
## Days -0.138
```

3.7 Serial correlations in time and space

The hierarchical model of $y|x, u = x'\beta + z'u + \epsilon$ can work well for correlations within blocks, but not for correlations in time as the correlations decay in time. The following uses nlme package to calculate time serial data.

```
library(nlme)
head(nlme::Ovary,n=50)
```

```
## Grouped Data: follicles ~ Time | Mare
##      Mare      Time follicles
## 1      1 -0.13636360         20
## 2      1 -0.09090910         15
## 3      1 -0.04545455         19
## 4      1  0.00000000         16
## 5      1  0.04545455         13
## 6      1  0.09090910         10
## 7      1  0.13636360         12
## 8      1  0.18181820         14
## 9      1  0.22727270         13
## 10     1  0.27272730         20
## 11     1  0.31818180         22
## 12     1  0.36363640         15
## 13     1  0.40909090         18
## 14     1  0.45454550         17
## 15     1  0.50000000         14
## 16     1  0.54545450         18
## 17     1  0.59090910         14
## 18     1  0.63636360         16
## 19     1  0.68181820         17
## 20     1  0.72727270         18
## 21     1  0.77272730         18
## 22     1  0.81818180         17
## 23     1  0.86363640         14
## 24     1  0.90909090         12
## 25     1  0.95454550         12
## 26     1  1.00000000         14
## 27     1  1.04545500         10
## 28     1  1.09090900         11
## 29     1  1.13636400         16
## 30     2 -0.15000000          6
## 31     2 -0.10000000          6
## 32     2 -0.05000000          8
## 33     2  0.00000000          7
## 34     2  0.05000000         16
## 35     2  0.10000000         10
## 36     2  0.15000000         13
## 37     2  0.20000000          9
## 38     2  0.25000000          7
## 39     2  0.30000000          6
## 40     2  0.35000000          8
```

```
## 41    2  0.40000000    8
## 42    2  0.45000000    6
## 43    2  0.50000000    8
## 44    2  0.55000000    7
## 45    2  0.60000000    9
## 46    2  0.65000000    6
## 47    2  0.70000000    4
## 48    2  0.75000000    5
## 49    2  0.80000000    8
## 50    2  0.85000000   11
```

```
fm10var.lme <- nlme::lme(fixed=follicles ~ sin(2*pi*Time) + cos(2*pi*Time),
                        data = Ovary,
                        random = pdDiag(~sin(2*pi*Time)),
                        correlation=corAR1() )
summary(fm10var.lme)
```

```
## Linear mixed-effects model fit by REML
## Data: Ovary
##      AIC      BIC    logLik
## 1563.448 1589.49 -774.724
##
## Random effects:
## Formula: ~sin(2 * pi * Time) | Mare
## Structure: Diagonal
##      (Intercept) sin(2 * pi * Time) Residual
## StdDev:      2.858385          1.257977 3.507053
##
## Correlation Structure: AR(1)
## Formula: ~1 | Mare
## Parameter estimate(s):
##      Phi
## 0.5721866
## Fixed effects: follicles ~ sin(2 * pi * Time) + cos(2 * pi * Time)
##
##      Value Std.Error DF   t-value p-value
## (Intercept)    12.188089 0.9436602 295 12.915760 0.0000
## sin(2 * pi * Time) -2.985297 0.6055968 295 -4.929513 0.0000
## cos(2 * pi * Time) -0.877762 0.4777821 295 -1.837159 0.0672
## Correlation:
##      (Intr) s(*p*T
## sin(2 * pi * Time) 0.000
## cos(2 * pi * Time) -0.123 0.000
##
## Standardized Within-Group Residuals:
##      Min      Q1      Med      Q3      Max
## -2.34910093 -0.58969626 -0.04577893 0.52931186 3.37167486
```

```
##  
## Number of Observations: 308  
## Number of Groups: 11
```


Chapter 4

Basic Stat Concepts

4.1 Score

The score is the gradient (the vector of partial derivatives) of $\log L(\theta)$, with respect to an m -dimensional parameter vector θ .

$$S(\theta) = \frac{\partial \ell}{\partial \theta}$$

Typically, they use ∇ to denote the partial derivative.

$$\nabla \ell$$

Such differentiation will generate a $m \times 1$ row vector, which indicates the sensitivity of the likelihood.

Quote from Steffen Lauritzen's slides: "Generally the solution to this equation must be calculated by iterative methods. One of the most common methods is the Newton–Raphson method and this is based on successive approximations to the solution, using Taylor's theorem to approximate the equation."

For instance, using logit link, we can get the first derivative of log likelihood logistic regression as follows. We can not really find β easily to make the equation to be 0.

$$\begin{aligned} \frac{\partial \ell}{\partial \beta} &= \sum_{i=1}^n x_i^T \left[y_i - \frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}} \right] \\ &= \sum_{i=1}^n x_i^T [y_i - \hat{y}_i] \end{aligned}$$

4.2 Gradient and Jacobian

Remarks: This part discusses gradient in a more general sense.

When $f(x)$ is only in a single dimension space:

$$\mathbb{R}^n \rightarrow \mathbb{R}$$

$$\nabla f(x) = \left[\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \dots, \frac{\partial f}{\partial x_n} \right]$$

When $f(x)$ is only in a m-dimension space (i.e., Jacobian): $\mathbb{R}^n \rightarrow \mathbb{R}^m$

$$Jac(f) = \begin{bmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \frac{\partial f_1}{\partial x_3} & \dots & \frac{\partial f_1}{\partial x_n} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \frac{\partial f_2}{\partial x_3} & \dots & \frac{\partial f_2}{\partial x_n} \\ \dots & \dots & \dots & \dots & \dots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \frac{\partial f_m}{\partial x_3} & \dots & \frac{\partial f_m}{\partial x_n} \end{bmatrix}$$

For instance,

$$\mathbb{R}^n \rightarrow \mathbb{R}:$$

$$f(x, y) = x^2 + 2y$$

$$\nabla f(x, y) = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right] = [2x, 2]$$

$$\mathbb{R}^n \rightarrow \mathbb{R}^m$$

$$f(x, y) = (x^2 + 2y, x^3)$$

$$Jac(f) = \begin{bmatrix} 2x & 2 \\ 3x^2 & 0 \end{bmatrix}$$

4.3 Hessian and Fisher Information

Hessian matrix or Hessian is a square matrix of second-order partial derivatives of a scalar-valued function, or scalar field.

$$\mathbb{R}^n \rightarrow \mathbb{R}$$

$$Hessian = \nabla^2(f) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \frac{\partial^2 f}{\partial x_1 \partial x_3} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \frac{\partial^2 f}{\partial x_2 \partial x_3} & \dots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \frac{\partial^2 f}{\partial x_3 \partial x_1} & \frac{\partial^2 f}{\partial x_3 \partial x_2} & \frac{\partial^2 f}{\partial x_3^2} & \dots & \frac{\partial^2 f}{\partial x_3 \partial x_n} \\ \dots & \dots & \dots & \dots & \dots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \frac{\partial^2 f}{\partial x_n \partial x_3} & \dots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

As a special case, in the context of logit:

Suppose that the log likelihood function is $\ell(\theta)$. θ is a m dimension vector.

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \\ \theta_4 \\ \dots \\ \theta_m \end{bmatrix}$$

$$Hessian = \nabla^2(\ell) = \begin{bmatrix} \frac{\partial^2 \ell}{\partial \theta_1^2} & \frac{\partial^2 \ell}{\partial \theta_1 \partial \theta_2} & \frac{\partial^2 \ell}{\partial \theta_1 \partial \theta_3} & \dots & \frac{\partial^2 \ell}{\partial \theta_1 \partial \theta_m} \\ \frac{\partial^2 \ell}{\partial \theta_2 \partial \theta_1} & \frac{\partial^2 \ell}{\partial \theta_2^2} & \frac{\partial^2 \ell}{\partial \theta_2 \partial \theta_3} & \dots & \frac{\partial^2 \ell}{\partial \theta_2 \partial \theta_m} \\ \frac{\partial^2 \ell}{\partial \theta_3 \partial \theta_1} & \frac{\partial^2 \ell}{\partial \theta_3 \partial \theta_2} & \frac{\partial^2 \ell}{\partial \theta_3^2} & \dots & \frac{\partial^2 \ell}{\partial \theta_3 \partial \theta_m} \\ \dots & \dots & \dots & \dots & \dots \\ \frac{\partial^2 \ell}{\partial \theta_m \partial \theta_1} & \frac{\partial^2 \ell}{\partial \theta_m \partial \theta_2} & \frac{\partial^2 \ell}{\partial \theta_m \partial \theta_3} & \dots & \frac{\partial^2 \ell}{\partial \theta_m \partial \theta_m} \end{bmatrix}$$

“In statistics, the observed information, or observed Fisher information, is the negative of the second derivative (the Hessian matrix) of the “log-likelihood” (the logarithm of the likelihood function). It is a sample-based version of the Fisher information.” (Direct quote from Wikipedia.)

Thus, the observed information matrix:

$$-Hessian = -\nabla^2(\ell)$$

Expected (Fisher) information matrix:

$$E[-\nabla^2(\ell)]$$

4.4 Canonical link function

Inspired by a Stack Exchange post, I created the following figure:

$$\frac{Paramter}{\theta} \longrightarrow \gamma'(\theta) = \mu \longrightarrow \frac{Mean}{\mu} \longrightarrow g(\mu) = \eta \longrightarrow \frac{Linearpredictor}{\eta}$$

For the case of n time Bernoulli (i.e., Binomial), its canonical link function is logit. Specifically,

$$\frac{\text{Parameter}}{\theta = \beta^T x_i} \rightarrow \gamma'(\theta) = \frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}} \rightarrow \frac{\text{Mean}}{\mu = \frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}}} \rightarrow g(\mu) = \log \frac{\frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}}}{1 - \frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}}} \rightarrow \frac{\text{Linear predictor}}{\eta = \beta^T x_i}$$

Thus, we can see that,

$$\theta \equiv \eta$$

The link function $g(\mu)$ relates the linear predictor $\eta = \beta^T x_i$ to the mean μ .

Remarks:

- (1) Parameter is $\theta = \beta^T x_i$ (Not μ !).
- (2) $\mu = p(y = 1) = \frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}}$ (Not logit!).
- (3) Link function (i.e., $g(\mu)$) = logit = logarithm of odds = $\log \frac{\text{Event-Happened}}{\text{Event-Not-Happened}}$.
- (4) $g(\mu) = \log \frac{\mu}{1-\mu} = \beta^T x_i$. Thus, link function = linear predictor = log odds!
- (5) Quote from the Stack Exchange post “Newton Method and Fisher scoring for finding the ML estimator coincide, these links simplify the derivation of the MLE.”

(Recall, we know that μ or $p(y = 1)$ is the mean function. Recall that, n trials of coin flips, and get p heads. Thus $\mu = \frac{p}{n}$.)

4.5 Ordinary Least Squares (OLS)

Suppose we have n observation, and m variables.

$$\begin{bmatrix} x_{11} & x_{12} & x_{13} & \dots & x_{1m} \\ x_{21} & x_{22} & x_{23} & \dots & x_{2m} \\ \dots & & & & \\ x_{n1} & x_{n2} & x_{n3} & \dots & x_{nm} \end{bmatrix}$$

Thus, we can write it as the following n equations.

$$y_1 = \beta_0 + \beta_1 x_{11} + \beta_2 x_{12} + \dots + \beta_m x_{1m}$$

$$y_2 = \beta_0 + \beta_1 x_{21} + \beta_2 x_{22} + \dots + \beta_m x_{2m}$$

$$y_3 = \beta_0 + \beta_1 x_{31} + \beta_2 x_{32} + \dots + \beta_m x_{3m}$$

...

$$y_n = \beta_0 + \beta_1 x_{n1} + \beta_2 x_{n2} + \dots + \beta_m x_{nm}$$

We can combine all the n equations as the following one:

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_m x_{im} (i \in [1, n])$$

We can further rewrite it as a matrix format as follows.

$$y = X\beta$$

Where,

$$y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \\ \dots \\ y_n \end{bmatrix}$$

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} & x_{13} & \dots & x_{1m} \\ 1 & x_{21} & x_{22} & x_{23} & \dots & x_{2m} \\ \dots & & & & & \\ 1 & x_{n1} & x_{n2} & x_{n3} & \dots & x_{nm} \end{bmatrix}$$

$$\beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \beta_3 \\ \dots \\ \beta_m \end{bmatrix}$$

Since later we need the inverse of X , we need to make it into a square matrix.

$$X^T y = X^T X \hat{\beta} \Rightarrow \hat{\beta} = (X^T X)^{-1} X^T y$$

We can use R to implement this calculation. As we can see, there is no need to do any iterations at all, but rather just pure matrix calculation.

```
X<-matrix(rnorm(1000),ncol=2) # we define a 2 column matrix, with 500 rows
X<-cbind(1,X) # add a 1 constant
beta_true<-c(2,1,2) # True regression coefficients
beta_true<-as.matrix(beta_true)
y=X%%beta_true+rnorm(500)
```

```
transposed_X<-t(X)
beta_hat<-solve(transposed_X%*%X)%*%transposed_X%*%y
beta_hat
```

```
##           [,1]
## [1,] 2.017690
## [2,] 1.054682
## [3,] 2.037671
```

Side Notes The function of `as.matrix` will automatically make `c(2,1,2)` become the dimension of 3×1 , you do not need to transpose the β .

4.6 Taylor series

$$\begin{aligned} f(x)|_a &= f(a) + \frac{f'(a)}{1!}(x-a) + \frac{f''(a)}{2!}(x-a)^2 + \frac{f'''(a)}{3!}(x-a)^3 + \dots \\ &= \sum_{n=0}^{\infty} \frac{f^n(a)}{n!}(x-a)^n \end{aligned}$$

For example:

$$\begin{aligned} e^x|_{a=0} &= e^a + \frac{e^a}{1!}(x-a) + \frac{e^a}{2!}(x-a)^2 + \dots + \frac{e^a}{n!}(x-a)^n \\ &= 1 + \frac{1}{1!}x + \frac{1}{2!}x^2 + \dots + \frac{1}{n!}x^n \end{aligned}$$

if $x = 2$

$$e^2 = 7.389056$$

$$e^2 \approx 1 + \frac{1}{1!}x = 1 + \frac{1}{1!}2 = 3$$

$$e^2 \approx 1 + \frac{1}{1!}x + \frac{1}{2!}x^2 = 1 + \frac{1}{1!}2 + \frac{1}{2!}2 = 5 \dots$$

$$e^2 \approx 1 + \frac{1}{1!}x + \frac{1}{2!}x^2 + \frac{1}{3!}x^2 + \frac{1}{4!}x^2 + \frac{1}{5!}x^2 = 7.2666\dots$$

4.7 Fisher scoring

[I will come back to this later.]

<https://www2.stat.duke.edu/courses/Fall00/sta216/handouts/diagnostics.pdf>

<https://stats.stackexchange.com/questions/176351/implement-fisher-scoring-for-linear-regression>

4.8 References

1. Steffen Lauritzen's slides:
<http://www.stats.ox.ac.uk/~steffen/teaching/bs2HT9/scoring.pdf>
2. The Stack Exchange post:
<https://stats.stackexchange.com/questions/40876/what-is-the-difference-between-a-link-function-and-a-canonical-link-function>
3. Wikipedia for OLS
https://en.wikipedia.org/wiki/Ordinary_least_squares
4. Gradient and Jacobian
<https://math.stackexchange.com/questions/1519367/difference-between-gradient-and-jacobian>
https://www.youtube.com/watch?v=3xVMVT-2_t4
<https://math.stackexchange.com/questions/661195/what-is-the-difference-between-the-gradient-and-the-directional-derivative>
5. Hessian
https://en.wikipedia.org/wiki/Hessian_matrix
6. Observed information
https://en.wikipedia.org/wiki/Observed_information
7. Fisher information
https://people.missouristate.edu/songfengzheng/Teaching/MTH541/Lecture%20notes/Fisher_info.pdf
8. Link function
https://en.wikipedia.org/wiki/Generalized_linear_model#Link_function
<https://stats.stackexchange.com/questions/40876/what-is-the-difference-between-a-link-function-and-a-canonical-link-function>

Chapter 5

Basic R

This section is about R coding.

5.1 apply, lapply, sapply

5.1.1 apply

```
m_trying <- matrix(C<-(1:10),nrow=2, ncol=5)
m_trying
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    3    5    7    9
## [2,]    2    4    6    8   10
```

```
## Operating on the columns
apply(m_trying, 2, sum)
```

```
## [1]  3  7 11 15 19
```

```
## Operating on the rows
apply(m_trying, 1, sum)
```

```
## [1] 25 30
```

5.1.2 lapply

“lapply returns a list of the same length as X, each element of which is the result of applying FUN to the corresponding element of X.”

lapply operates on lists. Thus, as we can see below, even if m_trying is not a list, each cell becomes a list.

```
results1<-lapply(m_trying,sum)
str(results1)
```

```
## List of 10
## $ : int 1
## $ : int 2
## $ : int 3
## $ : int 4
## $ : int 5
## $ : int 6
## $ : int 7
## $ : int 8
## $ : int 9
## $ : int 10
```

```
is.list(results1)
```

```
## [1] TRUE
```

5.1.3 sapply

“sapply() function takes list, vector or data frame as input and gives output in vector or matrix.”

```
results2<-sapply(m_trying, sum)
str(results2)
```

```
## int [1:10] 1 2 3 4 5 6 7 8 9 10
```

```
is.list(results2)
```

```
## [1] FALSE
```

```
is.matrix(results2)
```

```
## [1] FALSE
```

```
is.data.frame(results2)
```

```
## [1] FALSE
```

```
is.vector(results2)
```

```
## [1] TRUE
```

5.2 C

```
mydata1<-matrix(runif(4*2),4,2)
mydata1
```

```
##           [,1]      [,2]
## [1,] 0.7767640 0.3839558
## [2,] 0.8404593 0.9506320
## [3,] 0.8705815 0.7041046
## [4,] 0.9530419 0.4219814
```

```
str(mydata1)
```

```
## num [1:4, 1:2] 0.777 0.84 0.871 0.953 0.384 ...
```

```
mydata2<-c(mydata1)
```

```
mydata2
```

```
## [1] 0.7767640 0.8404593 0.8705815 0.9530419 0.3839558 0.9506320 0.7041046
## [8] 0.4219814
```

```
str(mydata2)
```

```
## num [1:8] 0.777 0.84 0.871 0.953 0.384 ...
```


Chapter 6

Computing Techniques

Since GLMM can use EM algorithm in its maximum likelihood calculation (see McCulloch, 1994), it is practically useful to rehearse EM and other computing techniques.

6.1 Monte carlo approximation

Example: calculate the integral of $p(z > 2)$ when $z \sim N(0, 1)$. To use Monte Carlo approximation, we can have an indicator function, which will determine whether the sample from $N(0, 1)$ will be included into the calculation of the integral.

```
Nsim=10^4

indicator=function(x){
  y=ifelse((x>2),1,0)
  return(y)}


newdata<-rnorm(Nsim, 0,1 )

mc=c(); v=c(); upper=c(); lower=c()

for (j in 1:Nsim)
{
  mc[j]=mean(indicator(newdata[1:j]))
  v[j]=(j^-1)*var(indicator(newdata[1:j]))
  upper[j]=mc[j]+1.96*sqrt(v[j])
  lower[j]=mc[j]-1.96*sqrt(v[j])
}
```

```
library(ggplot2)
values=c(mc,upper,lower)
type=c(rep("mc",Nsim),rep("upper",Nsim),rep("lower",Nsim))
itr=rep(seq(1:Nsim),3)
data=data.frame(val=values, tp=type, itr=itr)
Rcode<-ggplot(data,aes(itr,val,col=tp))+geom_line(size=0.5)
Rcode+geom_hline(yintercept=1-pnorm(2),color="green",size=0.5)
```

Warning: Removed 2 rows containing missing values (geom_path).



6.2 Importance sampling

Importance sampling has samples generated from a different distribution than the distribution of interest. Specifically, assume that we want to calculate the expected value of $h(x)$, and $x \sim f(x)$.

$$E(h(x)) = \int h(x)f(x)dx = \int h(x)\frac{f(x)}{g(x)}g(x)dx$$

We can sample x_i from $g(x)$ and then calculate the mean of $h(x_i)\frac{f(x_i)}{g(x_i)}$.

Using the same explain above, we can use a shifted exponential distribution to help calculate the integral for normal distribution. Specifically,


$$\int_2^{\infty} \frac{1}{2\pi} e^{-\frac{1}{2}x^2} dx = \int_2^{\infty} \frac{\frac{1}{2\pi} e^{-\frac{1}{2}x^2}}{e^{-(x-2)}} e^{-(x-2)} dx$$

The idea is that, we can generate x_i from exponential distribution of $e^{-(x-2)}$, and then insert them into the targeted “expected (value) function” of $\frac{\frac{1}{2\pi} e^{-\frac{1}{2}x^2}}{e^{-(x-2)}}$. Thus, as you can see, importance sampling is based on the law of large numbers (i.e., If the same experiment or study is repeated independently a large number of times, the average of the results of the trials must be close to the expected value). We can use it to calculate integral based on link of the definition of expected value.

```
Nsim=10^4
normal_density=function(x)
{y=(1/sqrt(2*pi))*exp(-0.5*(x^2))
return(y)}
x=2-log(runif(Nsim))
ImpS=c(); v=c(); upper=c(); lower=c()
for (j in 1:Nsim)
{
ImpS[j]=mean(normal_density(x[1:j])/exp(-(x[1:j]-2)))
v[j]=(j^{-1})*var(normal_density(x[1:j])/exp(-(x[1:j]-2)))
upper[j]=ImpS[j]+1.96*sqrt(v[j])
lower[j]=ImpS[j]-1.96*sqrt(v[j])
}

library(ggplot2)
values=c(ImpS,upper,lower)
type=c(rep("mc",Nsim),rep("upper",Nsim),rep("lower",Nsim))
itr=rep(seq(1:Nsim),3)
data=data.frame(val=values, tp=type, itr=itr)
ggplot(data,aes(itr,val,col=tp))+geom_line(size=0.5)+
geom_hline(yintercept=1-pnorm(2),color="green",size=0.5)
```

```
## Warning: Removed 2 rows containing missing values (geom_path).
```



6.3 Newton Raphson algorithm

The main purpose of Newton Raphson algorithm is to calculate the root of a function (e.g., $x^2 - 3 = 0$). We know that in order to maximize the MLE, we need to calculate the first derivative of the function and then set it to zero $\ell'(x) = 0$. Thus, we can use the same Newton Raphson method to help calculate the MLE maximization as well.

There are different ways to understand Newton Raphson method, but I found the method fo geometric the most easy way to explain.

Specifically, suppose that you want to calculate the root of a function $f(x) = 0$. We assume the root is r . However, we do not that, and we randomly guess a point of a . Thus, we can get a tangent line with slope of $f'(a)$ and a point of $(a, f(a))$. Since we know the slope and one of its points, we can write the function for this tangent line.

$$y - f(a) = f'(a)(x - a)$$

To calculate the x -intercept, namely b in the figure, we can set $y = 0$, and get the following:

$$-f(a) = f'(a)(x - a) \Rightarrow x(or, b) = a - \frac{f(a)}{f'(a)}$$




Figure 6.1: Credit of this figure: <https://www.math.ubc.ca/~anstee/math104/newtonmethod.pdf>

If there is significant difference of $|a - b|$, we know that our original guess of a is not good. We better use b as the next guess, and calculate its tangent line again. To generalize, we can write it as follows.

$$x_{t+1} = x_t - \frac{f(x_t)}{f'(x_t)}$$

Okay, this method above is to calculate the root. For MLE, we can also use this method to calculate the root for the $\ell' = 0$. We can write it as follows.

$$x_{t+1} = x_t - \frac{\ell'(x_t)}{\ell''(x_t)}$$

Often, x is not just a single unknow parameter, but a vector. For this case, we can write it as follows.

$$\beta_{t+1} = \beta_t - \frac{\ell'(\beta_t)}{\ell''(\beta_t)}$$

6.3.1 Calculate the root

$$x^3 - 5 = 0$$


Note that, this is obviously not a maximization problem. In contrast, it involves a function with zero. As we can see, we can think it as the first order of Taylor approximation. That is, $f'(x) = x^3 - 5 = 0$. As we can see the following plot, it converts very quickly.

```

f_firstorder=function(x){x^3-5}
f_secondorder=function(x){3*x}
x_old=1;tolerance=1e-3
max_its=2000;iteration=1;difference=2
c_iteration<-c() ## to collect numbers generated in the iteration process
while(difference>tolerance & iteration<max_its){
  x_updated=x_old-(f_firstorder(x_old)/f_secondorder(x_old))
  difference=abs(x_updated-x_old);
  iteration=iteration+1;
  x_old=x_updated
  c_iteration<-c(c_iteration,x_updated)}

plot(c_iteration,type="b")

```



6.3.2 Logistic regression

Suppose we have n observation, and m variables.

$$\begin{bmatrix} x_{11} & x_{12} & x_{13} & \dots & x_{1m} \\ x_{21} & x_{22} & x_{23} & \dots & x_{2m} \\ \dots & & & & \\ x_{n1} & x_{n2} & x_{n3} & \dots & x_{nm} \end{bmatrix}$$

Typically, we add a vector of 1 being used to estimate the constant.

$$\begin{bmatrix} 1 & x_{11} & x_{12} & x_{13} & \dots & x_{1m} \\ 1 & x_{21} & x_{22} & x_{23} & \dots & x_{2m} \\ \dots & & & & & \\ 1 & x_{n1} & x_{n2} & x_{n3} & \dots & x_{nm} \end{bmatrix}$$

And, we have observe a vector of n y_i as well, which is a binary variable:

$$Y = \begin{bmatrix} 1 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ \dots \\ 1 \end{bmatrix}$$

Using the content from the MLE chapter, we can get:

$$\mathbf{L} = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{(1-y_i)}$$

Further, we can get a log-transformed format.

$$\log(\mathbf{L}) = \sum_{i=1}^n [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)]$$

Given that $p_i = \frac{e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}}{1 + e^{\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n}} = \frac{e^{\beta^T x}}{1 + e^{\beta^T x}}$, we can rewrite it as follows:

$$\log(\mathbf{L}) = \ell = \sum_{i=1}^n [y_i \log\left(\frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}}\right) + (1 - y_i) \log\left(1 - \frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}}\right)]$$

Before doing the derivative, we set.

$$\frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}} = p(\beta^T x_i)$$

$$\log(\mathbf{L}) = \ell = \sum_{i=1}^n [y_i \log(p(\beta^T x_i)) + (1 - y_i) \log(1 - p(\beta^T x_i))]$$

Note that, $\frac{\partial p(\beta^T x_i)}{\partial (\beta^T x_i)} = p(\beta^T x_i)(1 - p(\beta^T x_i))$. We will use it later.

$$\begin{aligned}
\nabla \ell &= \sum_{i=1}^n \left[y_i \frac{1}{p(\beta^T x_i)} \frac{\partial p(\beta^T x_i)}{\partial(\beta^T x_i)} \frac{\partial(\beta^T x_i)}{\partial \beta} + (1 - y_i) \frac{1}{1 - p(\beta^T x_i)} (-1) \frac{\partial p(\beta^T x_i)}{\partial(\beta^T x_i)} \frac{\partial(\beta^T x_i)}{\partial \beta} \right] \\
&= \sum_{i=1}^n x_i^T \left[y_i \frac{1}{p(\beta^T x_i)} p(\beta^T x_i) (1 - p(\beta^T x_i)) + (1 - y_i) \frac{1}{1 - p(\beta^T x_i)} (-1) p(\beta^T x_i) (1 - p(\beta^T x_i)) \right] \\
&= \sum_{i=1}^n x_i^T \left[y_i \frac{1}{p(\beta^T x_i)} p(\beta^T x_i) (1 - p(\beta^T x_i)) - (1 - y_i) \frac{1}{1 - p(\beta^T x_i)} p(\beta^T x_i) (1 - p(\beta^T x_i)) \right] \\
&= \sum_{i=1}^n x_i^T [y_i (1 - p(\beta^T x_i)) - (1 - y_i) p(\beta^T x_i)] \\
&= \sum_{i=1}^n x_i^T [y_i - y_i p(\beta^T x_i) - p(\beta^T x_i) + y_i p(\beta^T x_i)] \\
&= \sum_{i=1}^n x_i^T [y_i - p(\beta^T x_i)] \\
&= \sum_{i=1}^n x_i^T \left[y_i - \frac{e^{\beta^T x_i}}{1 + e^{\beta^T x_i}} \right]
\end{aligned}$$

As noted, the Newton Raphson algorithm needs the second order.

$$\begin{aligned}
\nabla^2 \ell &= \frac{\partial \sum_{i=1}^n x_i^T [y_i - p(\beta^T x_i)]}{\partial \beta} \\
&= - \sum_{i=1}^n x_i^T \frac{\partial p(\beta^T x_i)}{\partial \beta} \\
&= - \sum_{i=1}^n x_i^T \frac{\partial p(\beta^T x_i)}{\partial(\beta^T x_i)} \frac{\partial(\beta^T x_i)}{\partial \beta} \\
&= - \sum_{i=1}^n x_i^T p(\beta^T x_i) (1 - p(\beta^T x_i)) x_i
\end{aligned}$$

The following are the data simulation (3 IVs and 1 DV) and Newton Raphson analysis.

```

# Data generation
set.seed(123)
n=500
x1_norm<-rnorm(n)
x2_norm<-rnorm(n,3,4)
x3_norm<-rnorm(n,4,6)
x_combined<-cbind(1,x1_norm,x2_norm,x3_norm) # dimension: n*4
coefficients_new<-c(1,2,3,4) #true regression coefficient

```

```

inv_logit<-function(x,b){exp(x**b)/(1+exp(x**b))}
prob_generated<-inv_logit(x_combined,coefficients_new)
y<-c()
for (i in 1:n) {y[i]<-rbinom(1,1,prob_generated[i])}

# Newton Raphson

#We need to set random starting values.
beta_old<-c(1,1,1,1)
tolerance=1e-3
max_its=2000;iteration=1;difference=2
W<-matrix(0,n,n)

while(difference>tolerance & iteration<max_its)
{
  # The first order
  f_firstorder<-t(x_combined)**(y-inv_logit(x_combined,beta_old))
  # The second order
  diag(W) = inv_logit(x_combined,beta_old)*(1-inv_logit(x_combined,beta_old))
  f_secondorder<-t(x_combined)**W**x_combined
  # Calculate the beta_updated
  beta_updated=beta_old-(solve(f_secondorder)**f_firstorder)
  difference=max(abs(beta_updated-beta_old));
  iteration=iteration+1;
  beta_old=beta_updated}

beta_old

##           [,1]
##          0.9590207
## x1_norm 1.7974165
## x2_norm 3.0072303
## x3_norm 3.9578107

```

$$\begin{aligned}
\frac{\partial \ell}{\partial \beta} &= \sum_{i=1}^n \left[y_i \frac{1}{p(\beta^T x_i)} \frac{\partial p(\beta^T x_i)}{\partial (\beta^T x_i)} \frac{\partial (\beta^T x_i)}{\partial \beta} + (1-y_i) \frac{1}{1-p(\beta^T x_i)} (-1) \frac{\partial p(\beta^T x_i)}{\partial (\beta^T x_i)} \frac{\partial (\beta^T x_i)}{\partial \beta} \right] \\
&= \sum_{i=1}^n \left[y_i \frac{1}{p(\beta^T x_i)} \phi(\beta^T x_i) - (1-y_i) \frac{1}{1-p(\beta^T x_i)} \phi(\beta^T x_i) \right] x_i
\end{aligned}$$

$$\Phi(\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3) = p(y = 1)$$

```

# Data generation
n=500
x1_norm<-rnorm(n)
x2_norm<-rnorm(n)
x3_norm<-rnorm(n)
x_combined<-cbind(1,x1_norm,x2_norm,x3_norm)
coefficients_new<-c(2,2,3,3) #true regression coefficient
inv_norm<-function(x,b){pnorm(x*%b)}
prob_generated<-inv_norm(x_combined,coefficients_new)
y<-c()
for (i in 1:n) {y[i]<-rbinom(1,1,prob_generated[i])}

# Newton Raphson

#We need to set random starting values.
x_old<-c(1,1,1,1)
tolerance=1e-3
max_its=2000;iteration=1;difference=2

while(difference>tolerance & iteration<max_its){
  x_updated=x_old-(f_firstorder(x_old)/f_secondorder(x_old))
  difference=abs(x_updated-x_old);
  iteration=iteration+1;
  x_old=x_updated
  c_iteration<-c(c_iteration,x_updated)}

plot(c_iteration,type="b")

```

6.4 Metropolis Hastings

Metropolis–Hastings is a MCMC method for obtaining a sequence of random samples from a probability distribution from which direct sampling is difficult. By using the samples, we can plot the distribution (through histogram), or we can calculate the integral (e.g., you need to calculate the expected value).

(Side note: does this remind you the importance sampling? Very similar!)

Basic logic (my own summary):

- (1) Set up a random starting value of x_0 .
- (2) Sample a y_0 from the instrumental function of $q(x)$.
- (3) Calculate the following:

$$p = \frac{f(y_0) q(x_0)}{f(x_0) q(y_0)}$$

- (4) $\rho = \min(p, 1)$

$$(5) \ x_1 = \begin{cases} y_0 & p \\ x_0 & 1 - p \end{cases}$$

(6) Repeat n times (n is set subjectively.)

Use normal pdf to sample gamma distribution


```
alpha=2.7; beta=6.3 # I randomly chose alpha and beta values for the target gamma function

Nsim=5000 ## define the number of iteration

X=c(rgamma(1,1)) # initialize the chain from random starting numbers
mygamma<-function(Nsim,alpha,beta){
  for (i in 2:Nsim){
    Y=rnorm(1)
    rho=dgamma(Y,alpha,beta)*dnorm(X[i-1])/(dgamma(X[i-1],alpha,beta)*dnorm(Y))
    X[i]=X[i-1] + (Y-X[i-1])*(runif(1)<rho)
  }
  X
}

hist(mygamma(Nsim,alpha,beta), breaks = 100)
```

Histogram of mygamma(Nsim, alpha, beta)



6.5 EM

EM algorithm is an iterative method to find ML or maximum a posteriori (MAP) estimates of parameters.

Direct Ref: <http://www.di.fc.ul.pt/~jpn/r/EM/EM.html>

Suppose that we only observe X , and do not know Z . We thus need to construct the complete data of (X, Z) . Given $p(Z|X, \theta)$, we can compute the likelihood of the complete dataset:

$$p(X, Z|\theta) = p(Z|X, \theta)p(X|\theta)$$

The EM algorithm:

- (0) We got X and $p(Z|X, \theta)$
- (1) Random assign a θ_0 , since we do not know any of them.
- (2) E-step: $Q_{\theta_i} = E_{Z|X, \theta_i}[\log p(X, Z|\theta)]$
- (3) M-step: compute $\theta_{i+1} \leftarrow \operatorname{argmax} Q_{\theta_i}$
- (4) If θ_i and θ_{i+1} are not close enough, $\theta_i \leftarrow \theta_{i+1}$. Goto step 2.

For examples, you can refer to the following link: <http://www.di.fc.ul.pt/~jpn/r/EM/EM.html>

(It is `em_R.r` in `R_codes` folder. Personally, I can also refer to Quiz 2 in 536.)

6.6 References

1. The UBC PDF about Newton

<http://www.math.ubc.ca/~ansteemath104/newtonmethod.pdf>

2. Some other pages about Newton and logistic regression

<http://www.win-vector.com/blog/2011/09/the-simpler-derivation-of-logistic-regression/>

<https://stats.stackexchange.com/questions/344309/why-using-newtons-method-for-logistic-regression-often-fails>

<https://tomroth.com.au/logistic/>

<https://www.stat.cmu.edu/~cshalizi/350/lectures/26/lecture-26.pdf>

<https://www.stat.cmu.edu/~cshalizi/402/lectures/14-logistic-regression/lecture-14.pdf>

<http://hua-zhou.github.io/teaching/biostatm280-2017spring/slides/18-newton/newton.html>

3. MH

<https://www.youtube.com/watch?v=VGRVRjr0vyw>

Chapter 7

Generalized Linear Mixed Models

7.1 Basics of GLMM

Recall the formula in the probit model:

$$Y^* = X\beta + \epsilon, \epsilon \sim N(0, \sigma^2) = N(0, I)$$

Similar to LMM, binary model with random effect can be written as follows.

$$Y^* = X\beta + Zu + \epsilon$$

where,

$$\epsilon \sim N(0, I)$$

$$u \sim N(0, D)$$

We also assume ϵ and u are independent. Thus, we know that D represents the variances of the random effects. If we make $u = 1$, the model becomes the usual probit model. McCulloch (1994) states that there are a few advantages to use probit, rather than logit models. (Note that, however, probit is not canonical link function, but logit is!)

The following is the note from Charle E. McCulloch's "Maximum likelihood algorithms for Generalized Linear Mixed Models"

7.2 Some References

<http://www.biostat.umn.edu/~baolin/teaching/linmods/glmm.html>

http://www.biostat.umn.edu/~baolin/teaching/probmods/GLMM_mcmc.html

<https://bbolker.github.io/mixedmodels-misc/glmmFAQ.html>

Chapter 8

Twitter Example

The following is part of my course project for Stat 536. It aims to replicate part of the findings from Barbera (2015) Birds of the Same Feather Tweet Together: Bayesian Ideal Point Estimation Using Twitter Data. Political Analysis 23 (1). Note that, the following model is much simpler than that in the original paper.

8.1 Model

Suppose that a Twitter user is presented with a choice between following or not following another target $j \in \{1, \dots, m\}$. Let $y_j = 1$ if the user decides to follow j , and $y_j = 0$ otherwise.

$$y_j = \begin{cases} 1 & \text{Following} \\ 0 & \text{NotFollowing} \end{cases}$$

$$p(y_j = 1|\theta) = \frac{\exp(-\theta_0|\theta_1 - x_j|^2)}{1 + \exp(-\theta_0|\theta_1 - x_j|^2)}$$

We additionally know the priors of θ .

$$\theta_i \sim N(0, 10^2) (i = 0, 1)$$

The likelihood function is as follows.

$$L(Y|\theta) = \prod_{j=1}^m \left(\frac{\exp(-\theta_0|\theta_1 - x_j|^2)}{1 + \exp(-\theta_0|\theta_1 - x_j|^2)} \right)^{y_j} \left(1 - \frac{\exp(-\theta_0|\theta_1 - x_j|^2)}{1 + \exp(-\theta_0|\theta_1 - x_j|^2)} \right)^{(1-y_j)}$$

Thus, the posterior is as follows.

$$L(Y|\theta) \cdot N(\theta_0|0, 10) \cdot N(\theta_1|0, 10)$$

$$\propto \prod_{j=1}^m \left(\frac{\exp(-\theta_0|\theta_1 - x_j|^2)}{1 + \exp(-\theta_0|\theta_1 - x_j|^2)} \right)^{y_j} \left(1 - \frac{\exp(-\theta_0|\theta_1 - x_j|^2)}{1 + \exp(-\theta_0|\theta_1 - x_j|^2)} \right)^{(1-y_j)} \cdot \exp\left(-\frac{1}{2}\left(\frac{\theta_0}{10}\right)^2\right) \cdot \exp\left(-\frac{1}{2}\left(\frac{\theta_1}{10}\right)^2\right)$$

#Establish the function for logistic regression

```
Expit<-function(x){exp(x)/(1+exp(x))}
```

#Construct the posterior - in a log-format

#To make sure that the estimate of theta_1 is stable,

#the following code wants to make sure that theta_0 is always greater than zero.

```
log_post<-function(Y, X, theta)
{
  if(theta[1]<=0){post=-Inf}
  if(theta[1]>0){
    prob1<-Expit(-theta[1]*((theta[2]-X)^2))
    likelihood<-sum(dbinom(Y,1,prob1,log = TRUE))
    priors<-sum(dnorm(theta,0,10,log=TRUE))
    post=likelihood+priors}
  return(post)
}
```

```
Bayes_logit<-function (Y,X,n_samples=2000)
```

```
{
```

#Initial values

```
  theta<-c(5,5)
```

#store data

```
  keep.theta<-matrix(0,n_samples,2)
```

```
  keep.theta[1,]<-theta
```

#acceptance and rejection

```
  acc<-att<-rep(0,2)
```

#current log posterior

```
  current_lp<-log_post(Y,X,theta)
```

```
  for (i in 2:n_samples)
```

```
  {
```

```
    for(j in 1:2)
```

```
    {
```

#attempt + 1

```
      att[j]<-att[j]+1
```

```
      can_theta<-theta
```

```
      can_theta[j]<-rnorm(1,theta[j],0.5)
```

```

#candidate of log posterior
candidate_lp<-log_post(Y,X,can_theta)
Rho<-min(exp(candidate_lp-current_lp),1)
Random_probability<-runif(1)
if (Random_probability<Rho)
{
  theta<-can_theta
  current_lp<-candidate_lp
  #acceptance + 1, as long as Random_probability<Rho
  acc[j]<-acc[j]+1
}
}
#save theta
keep.theta[i,]<-theta
}
#Return: including theta and acceptance rate
list(theta=keep.theta,acceptance_rate=acc/att)
}

```

8.2 Simulating Data of Senators on Twitter

Assume that we have 100 senators, 50 Democrats and 50 Republicans, who we know their ideology. Assume that Democrats have negative ideology scores to indicate that they are more liberal, whereas Republicans have positive scores to indicate that they are more conservative. The following is data simulation for senators.

```

# Republicans are more conservative, and they have positive numbers.
Republicans<-c()
Republicans<-rnorm(50,1,0.5)
No_Republicans<-rep(1:50,1)
Part_1<-cbind(No_Republicans,Republicans)

# Democrats are more liberal, and they have negative numbers.
Democrats<-c()
Democrats<-rnorm(50,-1,0.5)
No_Democrats<-rep(51:100,1)
Part_2<-cbind(No_Democrats,Democrats)
Data_Elites<-rbind(Part_1,Part_2)
Data_Elites<-as.data.frame(Data_Elites)
colnames(Data_Elites) <- c("Elite_No","Elite_ideology")

head(Data_Elites)

##   Elite_No Elite_ideology

```

```
## 1      1      1.0541992
## 2      2      0.3805544
## 3      3      1.3568577
## 4      4      0.9922547
## 5      5      1.0089966
## 6      6      0.8878271
```

8.3 Simulating Data of Conservative Users on Twitter and Model Testing

Assume that we observe one Twitter user, who is more conservative. To simulate Twitter following data for this user, I assign this user to follow more Republican senators. Thus, if the Metropolis Hastings algorithm works as intended, we would expect to see a positive estimated value for their ideology. Importantly, as we can see in the histogram below, the estimated value indeed is positive, providing preliminary evidence for the statistical model and the algorithm. In addition, for the acceptance rate, we can see that the constant has a lower number than ideology, since we only accept a constant when it is positive.


```
#This user approximately follows 45 Republican Senators and 10 Democrat Senators.
Data_user<-as.data.frame(matrix(c(ifelse(runif(50)<.1,0,1),ifelse(runif(50)<.8,0,1))),
colnames(Data_user)<-c("R_User")
Data_combined<-cbind(Data_Elites,Data_user)
```

```
X_data<-Data_combined$Elite_ideology
Y_data<-Data_combined$R_User
```


```
fit_C<-Bayes_logit(Y_data,X_data)
fit_C$acceptance_rate
```

```
## [1] 0.1320660 0.5557779
```


```
plot(fit_C$theta[,1],main="Constant (Conservative Users)",
      xlab="Iteration Process",ylab="Estimated Scores",type="l")
```

Constant (Conservative Users)

```
plot(fit_C$theta[,2],main="Estimated Ideology Scores (Conservative Users)",
     xlab="Iteration Process",ylab="Ideology Scores",type="l")
```

Estimated Ideology Scores (Conservative Users)

```
hist(fit_C$theta[,2],main="Estimated Ideology Scores (Conservative Users)",
     xlab="Ideology Scores",breaks = 100)
```



8.4 Simulating Data of Liberal Users on Twitter and Model Testing

To further verify the Metropolis Hastings algorithm, I plan to test the opposite estimate. Specifically, assume that we observe another user, who is more liberal. To simulate Twitter following data for this user, I assign this user to follow more Democrat senators. In this case, we would expect to see a negative value for their estimated ideology. As we can see in the histogram shown below, as expected, the estimated value is negative, providing convergent evidence for the model and the algorithm.


```
#This user approximately follows 10 Republican Senators and 45 Democrat Senators.
Data_user<-as.data.frame(matrix(c(ifelse(runif(50)<.8,0,1),ifelse(runif(50)<.1,0,1))),
colnames(Data_user)<-c("L_User")
Data_combined<-cbind(Data_Elites,Data_user)

X_data<-Data_combined$Elite_ideology
Y_data<-Data_combined$L_User


fit_L<-Bayes_logit(Y_data,X_data)
fit_L$acceptance_rate
```


8.4. SIMULATING DATA OF LIBERAL USERS ON TWITTER AND MODEL TESTING73


```
## [1] 0.1585793 0.5092546  
plot(fit_L$theta[,1],main="Constant (Liberal Users)",  
      xlab="Iteration Process",ylab="Estimated Scores",type="l")
```



```
plot(fit_L$theta[,2],main="Estimated Ideology Scores (Liberal Users)",  
      xlab="Iteration Process",ylab="Ideology Scores",type="l")
```



```
hist(fit_L$theta[,2],main="Estimated Ideology Scores (Liberal Users)",  
     xlab="Ideology Scores",breaks = 100)
```



Chapter 9

Practice: Learning on the Battle Field

9.1 R code

```
#https://fivethirtyeight.com/contributors/josh-hermsmeyer/
# https://github.com/ryurko/nflscrapR-data/blob/master/legacy_data/README.md

mydata1 = read.csv('plays.txt')
unique(mydata1$gameId)

unique(mydata1$PassLength)
table(mydata1$PassLength)
table(mydata1$PassResult)
table(mydata1$numberOfPassRushers)

##mydata3 = read.csv(url('https://raw.githubusercontent.com/ryurko/nflscrapR-data/master/legacy_data/2017playbyplay.csv'))
##write.csv(mydata3, '2017playbyplay.csv')

mydata3<-read.csv('2017playbyplay.csv')
nrow(mydata3)
table(mydata3$Passer)
table(mydata3$PlayType)

mydata5<-mydata3[!duplicated(mydata3[,c('GameID', 'Passer')]),]
unique(mydata3$GameID)
mydata6<-subset(mydata3,down==1)
```

```

mydata7<-subset(mydata6,PlayType=='Pass'|PlayType=='Run')
#table(mydata7$PlayType)
#table(droplevels(mydata7$PlayType))

mydata7$PlayType<-droplevels(mydata7$PlayType)
table(mydata7$PlayType)

#http://rstudio-pubs-static.s3.amazonaws.com/6975_c4943349b6174f448104a5513fed59a9.htm
source("http://pcwww.liv.ac.uk/~william/R/crosstab.r")
mydata8<-mydata7[,c('Passer','PlayType','GameID','posteam','DefensiveTeam','Yards.Gain')]
#results<-crosstab(mydata8, row.vars = "GameID", col.vars = "PlayType", type = "r")
#p1<-results$crosstab
#hist(p1[,1],20)

library(plyr)
count_vector<-count(mydata8, "GameID")

l_new<-length(count_vector$freq)
time<-c()
for(i in 1:l_new)
{time<-append(time,rep(1:count_vector$freq[i]))}
nrow(time)
mydata8$time<-time
mydata8$play_new<-ifelse(mydata8$PlayType=='Pass',1,0)

n_counting<-0 # help counting the number of pairs

## The following code collects all the rows of each pair. However, it is difficult to
# in such a format.

#empty_df = mydata8[FALSE,]
#for (i in 1:l_new) # level of different game
#{
#   for(j in 1:((count_vector$freq[i])-1)) # within the same game
#   {
#       if(i==1)
#       {row_id<-j}
#       else {row_id<-sum(count_vector$freq[1:(i-1)])+j}
#
#       #print(row_id)
#       if(as.character(mydata8[row_id,]$posteam)!=as.character(mydata8[row_id+1,]$posteam))
#       {
#           print("not same team")
#       }
#   }
#}

```

```

#       if (nrow(empty_df)==0)
#         {empty_df<-mydata8[row_id:(row_id+1),]}
#       else
#         {
#           if(row.names(mydata8[row_id,])!=row.names(tail(empty_df,1)))
#             {empty_df<-rbind(empty_df,mydata8[row_id,])}
#           empty_df<-rbind(empty_df,mydata8[row_id+1,])
#         }
#       n_counting<-n_counting+1
#     }
#   }
#}

# The following code only collects the second row of the pair, but adds data of
### PT_L: type of play in the last first down from the other team
### TG_L: Yards.Gained in the last play
### FirstDown: did they get first down or not. Note that, if yes, it means it was a fumble.

PT_L="Pass"
TG_L=0
FD_L=0

pari_data= mydata8[1,]
pari_data<-cbind(pari_data,PT_L,TG_L,FD_L)
pari_data<-pari_data[FALSE,]

for (i in 1:l_new) # level of different game
{
  for(j in 1:((count_vector$freq[i])-1)) # within the same game
  {

    if(i==1)
    {row_id<-j}
    else {row_id<-sum(count_vector$freq[1:(i-1)])+j}

    print(row_id)
    if(as.character(mydata8[row_id,]$posteam)!=as.character(mydata8[row_id+1,]$posteam))
    {
      print("not same team")
      PT_L<-as.character(mydata8[row_id,]$PlayType)
      TG_L<-mydata8[row_id,]$Yards.Gained
      FD_L<-mydata8[row_id,]$FirstDown

      new_row<-cbind(mydata8[(row_id+1),],PT_L,TG_L,FD_L)

```

```

    pari_data<-rbind(pari_data,new_row)
  }

  n_counting<-n_counting+1
}
}

pari_data$same<-ifelse(pari_data$PlayType==pari_data$PT_L,1,0)

#write.csv(pari_data, 'pari_data.csv')

write.table(pari_data, file = "pari_data.csv",row.names=FALSE,na = "", sep=",")

```

Remarks

1. mylogit1: in general, a team has a different play in their first down, compared to the other team in the last first down.
2. mylogit2: If the defence team passed in the last first down, the offence team is less likely to use pass. If the defence team gained more yards, the offence team is more likely to pass in the next first down. If the defence team fumbled, it will reduce the chance the offence team to do the pass.

```

pari_data2<-read.csv('pari_data.csv')

mylogit1 = glm(same~1, family=binomial, data=pari_data2)
summary(mylogit1)

##
## Call:
## glm(formula = same ~ 1, family = binomial, data = pari_data2)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.117   -1.117   -1.117    1.239    1.239
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -0.14395    0.02809  -5.124   3e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 7035.5  on 5093  degrees of freedom
## Residual deviance: 7035.5  on 5093  degrees of freedom
## AIC: 7037.5

```

```
##
## Number of Fisher Scoring iterations: 3
mylogit2 = glm(play_new~same+TG_L+FD_L, family=binomial, data=pari_data2)
summary(mylogit2)

##
## Call:
## glm(formula = play_new ~ same + TG_L + FD_L, family = binomial,
##      data = pari_data2)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6114  -0.9783  -0.9382   1.0995   1.5672
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.175629   0.040712   4.314 1.6e-05 ***
## same        -0.757822   0.057618 -13.152 < 2e-16 ***
## TG_L         0.010439   0.003873   2.695 0.00704 **
## FD_L        -0.268115   0.148835  -1.801 0.07164 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 7034.3  on 5093  degrees of freedom
## Residual deviance: 6850.1  on 5090  degrees of freedom
## AIC: 6858.1
##
## Number of Fisher Scoring iterations: 4
library(lme4)
mylogit3 = glmer(same~play_new+TG_L+FD_L+(1|GameID), family= binomial("logit"), data=pari_data2)

## boundary (singular) fit: see ?isSingular
summary(mylogit3)

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: same ~ play_new + TG_L + FD_L + (1 | GameID)
## Data: pari_data2
##
##      AIC      BIC    logLik deviance df.resid
## 6862.4  6895.1 -3426.2  6852.4     5089
##
```

```
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -1.3918 -0.7763 -0.7532  0.9061  1.6255
##
## Random effects:
##   Groups Name            Variance Std.Dev.
##   GameID (Intercept) 1.562e-15 3.953e-08
## Number of obs: 5094, groups:  GameID, 256
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.197140   0.040513   4.866 1.14e-06 ***
## play_new     -0.757838   0.057619 -13.153 < 2e-16 ***
## TG_L         0.006027   0.003824   1.576 0.11502
## FD_L        -0.392792   0.150715  -2.606 0.00916 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Correlation of Fixed Effects:
##              (Intr) ply_nw TG_L
## play_new    -0.627
## TG_L        -0.270 -0.043
## FD_L        -0.147  0.031 -0.041
## convergence code: 0
## boundary (singular) fit: see ?isSingular

#Bill_1<- bild(play_new ~ TG_L+FD_L, data = mydata8, id="GameID",start = NULL, depende
#summary(Bill_1)

#locust2 <- bild(as.factor(PlayType) ~ time + I(time^2), data = mydata8,id="GameID",st
```

9.2 References

<https://arxiv.org/pdf/1403.7993.pdf>

http://www.dartmouth.edu/~chance/teaching_aids/books_articles/probability_book/Chapter11.pdf

<https://rpubs.com/JanpuHou/326048>

Chapter 10

Project Draft

```
mydata3<-read.csv('Schnibbe 1502 Binary Data.csv')
head(mydata3)
```

```
##   X0
## 1  0
## 2  1
## 3  0
## 4  0
## 5  1
## 6  0
```

```
NO_new<-rep(1:222)
mydata4<-cbind(mydata3,NO_new)
head(mydata4)
```

```
##   X0 NO_new
## 1  0      1
## 2  1      2
## 3  0      3
## 4  0      4
## 5  1      5
## 6  0      6
```

```
a1 = glmer(X0 ~ 1 + (1|NO_new), data = mydata4,family=binomial)
summary(a1)
```

```
## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: X0 ~ 1 + (1 | NO_new)
```

```
##      Data: mydata4
##
##      AIC      BIC    logLik deviance df.resid
##    243.3    250.1   -119.6    239.3     220
##
## Scaled residuals:
##      Min       1Q   Median       3Q      Max
## -0.5461 -0.5461 -0.5461 -0.5461  1.8311
##
## Random effects:
##  Groups Name      Variance Std.Dev.
## NO_new (Intercept) 1.246e-07 0.000353
## Number of obs: 222, groups: NO_new, 222
##
## Fixed effects:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.2098      0.1603   -7.549 4.38e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
a2 = glm(X0 ~ 1, data = mydata4, family=binomial)
summary(a2)
```

```
##
## Call:
## glm(formula = X0 ~ 1, family = binomial, data = mydata4)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7225  -0.7225  -0.7225  -0.7225   1.7151
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.2098      0.1595   -7.583 3.38e-14 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 239.29  on 221  degrees of freedom
## Residual deviance: 239.29  on 221  degrees of freedom
## AIC: 241.29
##
## Number of Fisher Scoring iterations: 4
```

10.1 Background

The following code is from this website: http://www.biostat.umn.edu/~baolin/teaching/probmods/GLMM_mcmc.html. I will remove it on this page after I complete my practice and learning.

In this example, it simulates a longitudinal data with 4 variables for each of 1000 separate individuals. Specifically, there are three continuous covariates (varying over time) and one ordinal covariate (constant over time). We will consider a random intercept model (mean zero and variance 100), and fit the data with `glmer()` from `lme4` R package.

The R code:

```
n = 1000; p = 3; K = 4; sig = 10
set.seed(123)

## time varying covariates
Xl = vector('list', K)
# 4 list, each 1000 individuals
for(i in 1:K) Xl[[i]] = matrix(rnorm(n*p), n,p)

## constant covariate
Z = rbinom(n, 2,0.2)

## random effects
#just 1000 random numubers?
U = rnorm(n)*sig

## fixed effects
# It ends a 1000*4 matrix
etaX = sapply(Xl, rowSums)

## random errors
eps = matrix(rnorm(n*K), n,K)

## logit model
eta = etaX + U + eps
# calculate probability
prb = 1/(1+exp(-eta))
D = 1*(matrix(runif(n*K),n,K)<prb) # comparing it to prb, and change to 1 and 0; 1000*4
# Select the first list from "Xl", and then add other 3 lists--> 4000 * 3
Xs = Xl[[1]]
for(k in 2:K) Xs = rbind(Xs, Xl[[k]])

## GLMM model
library(lme4)
```

```

sid = rep(1:n, K) # a vector of 1-1000, 4 repetitions
## model fit with GLMM (default to Laplace approximation)
# subjects as the random effect
a1 = glmer(c(D) ~ Xs + Z[sid] + (1|sid), family=binomial)

a1

```

```

## Generalized linear mixed model fit by maximum likelihood (Laplace
## Approximation) [glmerMod]
## Family: binomial ( logit )
## Formula: c(D) ~ Xs + Z[sid] + (1 | sid)
##      AIC      BIC    logLik deviance df.resid
## 3213.666 3251.430 -1600.833 3201.666     3994
## Random effects:
## Groups Name      Std.Dev.
## sid      (Intercept) 5.816
## Number of obs: 4000, groups: sid, 1000
## Fixed Effects:
## (Intercept)      Xs1      Xs2      Xs3      Z[sid]
##      0.1537      0.6650      0.6429      0.6074      0.0199

## MH sampling of random effects / data
## logit\Pr(D_i|eta_i,U) = eta_i+U; U \sim N(0,Vu)
## proposal dist: N(Uc,Vc)

```

```

U.mh <- function(Di,eta, Vu, Uc,Vc, B=100){
  ub = rep(0, B)
  ub[1] = rnorm(1)*sqrt(Vc)+Uc # random starting value
  prb = 1/(1+exp(-eta-ub[1]))
  llk0 = dnorm(ub[1],sd=sqrt(Vu), log=TRUE) + sum(log(Di*prb+(1-Di)*(1-prb))) - dnorm(
  for(k in 2:B){
    ub[k] = ub[k-1]
    uk = rnorm(1)*sqrt(Vc)+Uc
    prb = 1/(1+exp(-eta-uk))
    llk1 = dnorm(uk,sd=sqrt(Vu), log=TRUE) + sum(log(Di*prb+(1-Di)*(1-prb))) - dnorm(uk,
    alpha = exp( llk1 - llk0 )
    if(alpha>=1){
      ub[k] = uk
      llk0 = llk1
    } else{
      aa = runif(1)
      if(aa<alpha){
        ub[k] = uk
        llk0 = llk1
      }
    }
  }
}

```

```

    }
    return(ub)
  }

library(numDeriv)
UV.est <- function(Di,eta,Vu,Uc){
  llk0 = function(xpar){
    Uc = xpar
    prb = 1/(1+exp(-eta-Uc))
    res = dnorm(Uc,sd=sqrt(Vu), log=TRUE) + sum(log(Di*prb+(1-Di)*(1-prb)))
    -res
  }
  tmp = try(optim(Uc, llk0, method='Brent', lower=Uc-10,upper=Uc+10) )
  if(class(tmp)=='try-error') tmp = optim(Uc, llk0)
  Uc = tmp$par
  Vc = 1/hessian(llk0, Uc)
  c(Uc,Vc)
}

UV.mh <- function(Vu,beta,Uc, D,X,subj){
  ## Cov matrix
  sid = unique(subj); n = length(sid)
  Uc = Vc = rep(0,n)
  for(i in 1:n){
    ij = which(subj==sid[i]); ni = length(ij)
    Xi = X[ij,,drop=FALSE]
    eta = Xi%%beta
    zi = UV.est(D[ij],eta,Vu,Uc[i])
    Uc[i] = zi[1]; Vc[i] = zi[2]
  }
  return(list(Uc=Uc,Vc=Vc) )
}

#Newton Raphson update
# Compute first/second derives of complete data log likelihood
## score and fisher information
SF.mh <- function(Vu,beta,Uc,Vc, D,X,subj){
  ## S/hessian matrix
  sid = unique(subj); n = length(sid)
  p = dim(X)[2]
  S = rep(0, p)
  FI = matrix(0, p,p)
  sig2 = 0
  for(i in 1:n)
  {
    ij = which(subj==sid[i]); ni = length(ij)

```

```

    Xi = X[ij,,drop=FALSE]
    eta = Xi%%beta
    zi = U.mh(D[ij],eta,Vu,Uc[i],Vc[i], B=5e3)[- (1:1e3)]
    theta = sapply(eta, function(b0) mean(1/(1+exp(-b0-zi))) )
    theta2 = sapply(eta, function(b0) mean(exp(b0+zi)/(1+exp(b0+zi))^2) )
    FI = FI + t(Xi)%(theta2*Xi)
    S = S+colSums((D[ij]-theta)*Xi)
    sig2 = sig2 + mean(zi^2)
  }
  return(list(S=S, FI=FI, sig2=sig2/n) )
}

library(lme4)
sid = rep(1:n, K)
a1 = glmer(c(D) ~ Xs + Z[sid] + (1|sid), family=binomial)
## extract variance and fixed effects parameters; + mode/variance of (random effects/d
Vu = (getME(a1,'theta'))^2; beta = fixef(a1); Um = ranef(a1,condVar=TRUE)
D = c(D); X = cbind(1,Xs,Z[sid]); subj = sid
Uc = unlist(Um[[1]]); Vc = c( attr(Um[[1]], 'postVar') )
for(b in 1:100){
  ## NR updates with MH sampling
  obj = SF.mh(Vu,beta,Uc,Vc, D,X,subj)
  Vu = obj$sig2
  tmp = solve(obj$FI,obj$S)
  beta = beta + tmp
  ## Proposal dist update
  tmp1 = UV.mh(Vu,beta,Uc, D,X,subj)
  Uc = tmp1$Uc; Vc = tmp1$Vc
  cat(b, ': ', tmp, '; ', obj$S/n, '\n\t', sqrt(Vu), beta, '\n')
}

```

10.2 Important Examples with R code

1. Fitting mixed models with (temporal) correlations in R

https://bbolker.github.io/mixedmodels-misc/notes/corr_braindump.html

2. Mixed effects logistic regression

<https://stats.idre.ucla.edu/r/dae/mixed-effects-logistic-regression/>

10.3 References

1. Data

http://www.michelecoscia.com/?page_id=379

Chapter 11

Bayesian - 1

The following is the part of the class note that I took from the online course of “Bayesian Statistics: From Concept to Data Analysis.” (<https://www.coursera.org/learn/bayesian-statistics/home/welcome>)

Important note: All the notes here are just for my own study purpose. I do not claim any copyright. You can use it for study purpose as well, but not for any business purposes.

11.1 Frequentist perspective

$$\theta = \{fair, loaded\}$$

$$x \sim Bin(5, \theta)$$

$$\begin{aligned} f(x|\theta) &= \begin{cases} \binom{5}{x}(\frac{1}{2})^5 & \text{if } \theta = fair \\ \binom{5}{x}(0.7)^x(0.3)^{5-x} & \text{if } \theta = loaded \end{cases} \\ &= \binom{5}{x}(\frac{1}{2})^5 I_{\{\theta=fair\}} + \binom{5}{x}(0.7)^x(0.3)^{5-x} I_{\{\theta=loaded\}} \end{aligned}$$

When $x = 2$

$$f(\theta|x=2) = \begin{cases} \binom{5}{2}(\frac{1}{2})^5 = 0.3125 & \text{if } \theta = fair \\ \binom{5}{2}(0.7)^2(0.3)^{5-2} = 0.1323 & \text{if } \theta = loaded \end{cases}$$

Thus, based on MLE, it suggests that it should be “fair”, since it has a greater probability if we observe 2 head out of 5 trials.

However, we can not know the following probability: given that we observe $x = 2$, what is the probability that θ is fair?

$$P(\theta = fair|X = 2)$$

From the frequentist's perspective, the coin is the fixed coin. And thus, the probability of $P(\theta = fair|x = 2)$ is equal to $P(\theta = fair)$.

$$P(\theta = fair|x = 2) = P(\theta = fair)$$

As,

$$P(\theta = fair) \in C(0,1)(i.e., either 0 or 1)$$

11.2 Bayesian perspective

Prior $P(loaded) = 0.6$

$$\begin{aligned} f(\theta|X) &= \frac{f(x|\theta)f(\theta)}{\sum_{\theta} f(x|\theta)f(\theta)} \\ &= \frac{\binom{5}{x}[(\frac{1}{2})^5 \times 0.4 \times I_{\{\theta=fair\}} + (0.7)^x(0.3)^{5-x} \times 0.6 \times I_{\{\theta=loaded\}}]}{\binom{5}{x}[(\frac{1}{2})^5 \times 0.4 + (0.7)^x(0.3)^{5-x} \times 0.6]} \end{aligned}$$

$$\begin{aligned} f(\theta|X = 2) &= \frac{0.0125I_{\{\theta=fair\}} + 0.0079I_{\{\theta=loaded\}}}{0.0125 + 0.0079} \\ &= 0.612I_{\{\theta=fair\}} + 0.388I_{\{\theta=loaded\}} \end{aligned}$$

Thus, we can say that:

$$P(\theta = loaded|X = 2) = 0.388$$

We can change the prior, and get different posterior probabilities:

$$P(\theta = loaded) = \frac{1}{2} \rightarrow P(\theta = loaded|X = 2) = 0.297$$

$$P(\theta = loaded) = \frac{9}{10} \rightarrow P(\theta = loaded|X = 2) = 0.792$$

11.3 Continuous parameters

In the examples above, θ is discrete. In contrast, the examples below use continuous θ .

$$f(\theta|y) = \frac{f(y|\theta)f(\theta)}{f(y)} = \frac{f(y|\theta)f(\theta)}{\int f(y|\theta)f(\theta)d\theta} = \frac{\text{likelihood} \times \text{prior}}{\text{normalizing} - \text{constant}} \propto \text{likelihood} \times \text{prior}$$

Note that, the posterior is a PDF of θ , which is not in the function of $f(y)$. Thus, removing the denominator (i.e., the normalizing constant) does not change the form of the posterior.

11.3.1 Uniform

Suppose that θ is the probability of a coin getting head. We could assign a uniform distribution.

$$\theta \sim U[0, 1]$$

$$f(\theta) = I_{\{0 \leq \theta \leq 1\}}$$

(It is interesting to see how to write the pdf for uniform distribution.)

$$f(\theta|Y=1) = \frac{\theta^1(1-\theta)^0 I_{\{0 \leq \theta \leq 1\}}}{\int_{-\infty}^{+\infty} \theta^1(1-\theta)^0 I_{\{0 \leq \theta \leq 1\}} d\theta} = \frac{\theta I_{\{0 \leq \theta \leq 1\}}}{\int_0^1 \theta d\theta} = 2\theta I_{\{0 \leq \theta \leq 1\}}$$

If we ignore the normalizing constant, we will get

$$f(\theta|Y=1) \propto \theta^1(1-\theta)^0 I_{\{0 \leq \theta \leq 1\}} = \theta I_{\{0 \leq \theta \leq 1\}}$$

Thus, we can see that with vs. without the normalizing constant is the “2”.

11.3.2 Uniform: prior versus posterior

When θ follows uniform distribution:

Prior

$$P(0.025 < \theta < 0.975) = 0.95$$

$$P(0.05 < \theta) = 0.95$$

Posterior

$$P(0.025 < \theta < 0.975) = \int_{0.025}^{0.975} 2\theta d\theta = 0.95$$

$$P(0.05 < \theta) = 1 - P(\theta < 0.05) = \int_0^{0.05} 2\theta d\theta = 1 - 0.05^2 = 0.9975$$

Thus, we can see that, while $P(0.025 < \theta < 0.975)$ is the same for prior and posterior, $P(0.05 < \theta)$ is not the same.

11.3.3 Uniform: equal tailed versus HPD

Equal tailed

$$P(\theta < q|Y = 1) = \int_0^q 2\theta d\theta = q^2$$

$$P(\sqrt{0.025} < \theta < \sqrt{0.975}|Y = 1) = P(0.158 < \theta < 0.987) = 0.95$$

We can say that: there's a 95% probability that θ is in between 0.158 and 0.987.

Highest Posterior Density

$$P(\theta > \sqrt{0.05}|Y = 1) = P(\theta > 0.224|Y = 1) = 0.95$$

11.4 Bernoulli/binomial likelihood with uniform prior

$$\begin{aligned} f(\theta|Y = 1) &= \frac{\theta^{\sum y_i} (1 - \theta)^{n - \sum y_i} I_{\{0 \leq \theta \leq 1\}}}{\int_{-\infty}^{+\infty} \theta^{\sum y_i} (1 - \theta)^{n - \sum y_i} I_{\{0 \leq \theta \leq 1\}} d\theta} \\ &= \frac{\theta^{\sum y_i} (1 - \theta)^{n - \sum y_i} I_{\{0 \leq \theta \leq 1\}}}{\frac{\Gamma(\sum y_i + 1) \Gamma(n - \sum y_i + 1)}{\Gamma(n + 2)} \int_{-\infty}^{+\infty} \frac{\Gamma(n + 2)}{\Gamma(\sum y_i + 1) \Gamma(n - \sum y_i + 1)} \theta^{\sum y_i} (1 - \theta)^{n - \sum y_i} I_{\{0 \leq \theta \leq 1\}} d\theta} \\ &= \frac{\Gamma(n + 2)}{\Gamma(\sum y_i + 1) \Gamma(n - \sum y_i + 1)} \theta^{\sum y_i} (1 - \theta)^{n - \sum y_i} I_{\{0 \leq \theta \leq 1\}} \end{aligned}$$

Thus,

$$\theta|y \sim \text{Beta}(\sum y_i + 1, n - \sum y_i + 1)$$

Side note: $\text{Beta}(1, 1) = \text{Uniform}(0, 1)$:

$$\text{Beta}(\alpha, \beta) = \frac{x^{\alpha-1} (1-x)^{\beta-1}}{B(\alpha, \beta)} I_{\{0 \leq x \leq 1\}}$$

Thus, we can get the following since the support for beta distribution is $[0, 1]$:

$$\text{Beta}(1, 1) = \frac{x^0 (1-x)^0}{B(\alpha, \beta)} = 1 \times I_{\{0 \leq x \leq 1\}}$$

11.5 Conjugate priors

As noted above, beta prior (or, Uniform) leads to beta posterior. In a more general sense, Beta prior always leads to beta posterior.

For instance,

$$\begin{aligned}
 f(\theta|y) &\propto f(y|\theta)f(\theta) = \theta^{\sum y_i} (1-\theta)^{\sum n-y_i} \frac{\theta^{\alpha-1}(1-\theta)^{\beta-1}}{B(\alpha, \beta)} I_{\{0 \leq \theta \leq 1\}} \\
 &= \frac{1}{B(\alpha, \beta)} \theta^{\sum y_i + \alpha - 1} (1-\theta)^{\sum n - y_i + \beta - 1} I_{\{0 \leq \theta \leq 1\}} \\
 &\propto \theta^{\sum y_i + \alpha - 1} (1-\theta)^{\sum n - y_i + \beta - 1} I_{\{0 \leq \theta \leq 1\}}
 \end{aligned}$$

Thus,

$$f(\theta|y) \sim \text{Beta}(\alpha + \sum y_i, \beta + \sum n - y_i)$$

Conjugate prior: prior and posterior share the same distribution. As we can see, both the data and the prior contribute to the posterior.

For the prior of $\text{Beta}(\alpha, \beta)$, the mean is

$$\text{Mean}_{\text{prior}} = \frac{\alpha}{\alpha + \beta}$$

Posterior mean is,

$$\begin{aligned}
 \text{Mean}_{\text{posterior}} &= \frac{\alpha + \sum y_i}{\alpha + \sum y_i + \beta + n - \sum y_i} \\
 &= \frac{\alpha + \sum y_i}{\alpha + \beta + n} \\
 &= \frac{\alpha + \beta}{\alpha + \beta + n} \frac{\alpha}{\alpha + \beta} + \frac{n}{\alpha + \beta + n} \frac{\sum y_i}{n} \\
 &= \text{Weight}_{\text{prior}} \times \text{Mean}_{\text{prior}} + \text{Weight}_{\text{data}} \times \text{Mean}_{\text{data}}
 \end{aligned}$$

Side Note

- (1) Binomial proportion confidence interval:


$$\hat{p} \pm 1.96 \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$$

- (2) Mean of Beta distribution:


$$\frac{\alpha}{\alpha + \beta}$$

(3) Plot of Beta distribution

```
# Method 1  
x<-seq(0,1,length=200)  
plot(x,dbeta(x,1,5),type = "l")
```



```
# Method 2  
x2<-rbeta(200,1,5)  
hist(x2,prob = TRUE)  
lines(density(x2))
```



11.6 Poisson distribution

Pmf of Poisson distribution:

$$Pois(\lambda) \sim \frac{\lambda^k e^{-\lambda}}{k!}$$

We can replace k with the notation of y , and assume that we observe n y_i :

$$y_i \sim \frac{\lambda^{y_i} e^{-\lambda}}{y_i!}$$

$$f(y|\lambda) = \frac{\lambda^{\sum y_i} e^{-n\lambda}}{\prod_{i=1}^n y_i!}$$

We assume that λ follows Gamma distribution (i.e., Gamma prior):

$$\lambda \sim \Gamma(\alpha, \beta)$$

The pdf for Gamma distribution is:

$$\frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}$$

Thus, the posterior is as follows:

$$\begin{aligned}
 f(\lambda|y) &\propto f(y|\lambda)f(\lambda) \\
 &= \frac{\lambda^{\sum y_i} e^{-n\lambda}}{\prod_{i=1}^n y_i!} \times \frac{\beta^\alpha}{\Gamma(\alpha)} \lambda^{\alpha-1} e^{-\beta\lambda} \\
 &\propto \lambda^{\sum y_i} e^{-n\lambda} \times \lambda^{\alpha-1} e^{-\beta\lambda} \\
 &= \lambda^{(\alpha+\sum y_i)-1} e^{-(\beta+n)\lambda}
 \end{aligned}$$

Thus, the posterior is:

$$\Gamma(\alpha + \sum y_i, \beta + n)$$

As we know that, the mean of prior for Gamma is $\frac{\alpha}{\beta}$. Thus, we can get the mean for the posterior for Gamma is:

$$\begin{aligned}
 &= \frac{\alpha + \sum y_i}{\beta + n} \\
 &= \frac{\beta}{\beta + n} \frac{\alpha}{\beta} + \frac{n}{\beta + n} \frac{\sum y_i}{n}
 \end{aligned}$$

To determine the prior of α and β :

- (1) Prior mean $\frac{\alpha}{\beta}$
 - (a) Prior std. dev. $\frac{\sqrt{\alpha}}{\beta}$
 - (b) Effective sample size β
- (2) Vague prior Small $\varepsilon > 0$: $\Gamma(\varepsilon, \varepsilon)$. Thus, the posterior mean is primarily driven by the data:

$$\frac{\varepsilon + \sum y_i}{\varepsilon + n} \approx \frac{\sum y_i}{n}$$

- (1) As we know, beta prior lead the Bernoulli trial to a beta posterior. That is, we know $f(\theta|y) = \frac{f(y|\theta)f(\theta)}{f(y)}$. What is the prior predictive distribution of $f(y)$?
- (2) If Beta is Beta (3,3), what is the prior predictive probability that we will observe $y = 0$ in the next trial?

11.7 Exponential data

For instance, suppose that on average you need to wait for 10 minutes for a fast food delivery, and thus we can assume that $y \sim \text{Exp}(\lambda)$. Furthermore, we

assume that the prior λ follows Gamma distribution $\text{Gamma}(\alpha, \beta)$, thus it is with a mean of $\frac{\alpha}{\beta} = \frac{1}{10}$.

$$if \Gamma(100, 1000)$$

(Note that, it has a mean of $\frac{100}{1000} = \frac{1}{10}$).

Thus, we can get:

$$\begin{aligned} f(\lambda|y) &\propto f(y|\lambda)f(\lambda) \\ &\propto \lambda e^{-\lambda y} \lambda^{\alpha-1} e^{-\beta\lambda} \\ &\propto \lambda^{(\alpha+1)-1} e^{-(\beta+y)\lambda} \end{aligned}$$

Thus, we get

$$\lambda|y \sim \Gamma(\alpha + 1, \beta + y)$$

Thus, if we observe a data point that we need to wait for 12 minutes for a fast food delivery, we can update the posterior:

$$\lambda|y \sim \Gamma(101, 1012)$$

Thus, the mean for the posterior is

$$\frac{101}{1012} = \frac{1}{10.02}$$

Note:

- (1) Typically, we know that the pdf of Gamma is $\frac{\beta^\alpha}{\Gamma(\alpha)} x^{\alpha-1} e^{-\beta x}$. We replace x with λ since now the random variable of x is to represent the parameter λ in the exponential distribution.
- (2) In the above, we drop the constant part $(\frac{\beta^\alpha}{\Gamma(\alpha)})$ in the Gamma distribution as long as it does not include x (i.e., λ).
- (3) Suppose that you have 4 observations in total, then

$$\begin{aligned} f(\lambda|y) &\propto f(y_1|\lambda)f(y_2|\lambda)f(y_3|\lambda)f(y_4|\lambda)f(\lambda) \\ &\propto \lambda e^{-\lambda y_1} \lambda e^{-\lambda y_2} \lambda e^{-\lambda y_3} \lambda e^{-\lambda y_4} \lambda^{\alpha-1} e^{-\beta\lambda} \\ &\propto \lambda^{(\alpha+4)-1} e^{-(\beta+\sum_{i=1}^4 y_i)\lambda} \end{aligned}$$

Thus, the generalized form is as follows:

$$\lambda^{(\alpha+n)-1} e^{-(\beta+\sum_{i=1}^n y_i)\lambda}$$

11.8 Normal likelihood

11.8.1 When variance is known

$$x_i \sim N(\mu, \sigma_0^2)$$

(σ_0 is assumed to be known. Thus, the only unknown parameter is μ .)

The conjugate prior for normal distribution is normal distribution itself.

$$f(\mu|x) \sim f(x|\mu)f(\mu)$$

Assume that

$$\mu \sim N(m_0, s_0^2)$$

$$\mu|x \sim N\left(\frac{\frac{n\bar{x}}{\sigma_o^2} + \frac{m_o}{S_0^2}}{\frac{n}{\sigma_o^2} + \frac{1}{s_0^2}}, \frac{1}{\frac{n}{\sigma_o^2} + \frac{1}{s_0^2}}\right)$$

Thus, where

$$\begin{aligned} \frac{\frac{n\bar{x}}{\sigma_o^2} + \frac{m_o}{S_0^2}}{\frac{n}{\sigma_o^2} + \frac{1}{s_0^2}} &= \frac{\frac{n\bar{x}}{\sigma_o^2}}{\frac{n}{\sigma_o^2} + \frac{1}{s_0^2}} + \frac{\frac{m_o}{S_0^2}}{\frac{n}{\sigma_o^2} + \frac{1}{s_0^2}} \\ &= \frac{n}{n + \frac{\sigma_o^2}{s_0^2}} \bar{x} + \frac{\frac{\sigma_o^2}{S_0^2}}{n + \frac{\sigma_o^2}{s_0^2}} m_o \end{aligned}$$

Note:

- (1) As we can see, the posterior mean is a weighted mean – a combination of prior mean and sample mean.
- (2) When n is larger, the sample mean \bar{x} gets more weight.
- (3) For the prior mean m_0 , the smaller the prior variance s_0^2 is, the prior mean gets more weight. If the prior variance s_0^2 is big, the prior mean will get less weight in the final posterior mean.

11.8.2 When variance is unknown

$$x_i|\mu, \sigma^2 \sim N(\mu, \sigma^2)$$

$$\mu|\sigma^2 \sim N\left(m, \frac{\sigma^2}{w}\right)$$

Side note

$w = \frac{\sigma^2}{\sigma_\mu^2}$ effective sample size

$$\sigma^2 \sim \Gamma^{-1}(\alpha, \beta)$$

Thus, we can get that,

$$\sigma^2 | x \sim \Gamma^{-1}(\alpha + \frac{n}{2}, \beta + \frac{1}{2} \sum_{i=1}^n (x_i - \bar{x})^2 + \frac{nw}{2(n+w)} (\bar{x} - m)^2)$$

$$\mu | \sigma^2, x \sim N(\frac{n\bar{x} + wm}{n+w}, \frac{\sigma^2}{n+w})$$

Where,

$$\frac{n\bar{x} + wm}{n+w} = \frac{w}{n+w} m + \frac{n}{n+w} \bar{x}$$

$$\mu | x \sim t - \text{distribution}$$

11.9 Non-informative priors

11.9.1 Bernoulli

$$Y_i \sim B(\theta)$$

$$\theta \sim U[0, 1] = \text{Beta}(1, 1)$$

(Effective sample size is $1+1=2$)

If we get $\text{Beta}(\frac{1}{2}, \frac{1}{2})$ and $\text{Beta}(0.001, 0.001)$ have less impact on the posterior.

Improper prior, for instance, The prior

$$\text{Beta}(0, 0)$$

$$f(\theta) \propto \theta^{-1} (1 - \theta)^{-1}$$

In this case,

$$f(\theta | y) \propto \theta^{y-1} (1 - \theta)^{n-y-1} \sim \text{Beta}(y, n - y)$$

Posterior mean: $\frac{y}{n} = \hat{\theta}$

11.9.2 Gaussian

$$Y_i \sim N(\mu, \sigma^2)$$

Vague prior:

$$\mu \sim N(0, 1000000^2)$$

or,

$$f(\mu) \sim 1$$

$$\begin{aligned} f(\mu|y) &\propto f(y|\mu)f(\mu) \\ &\propto \exp\left(-\frac{1}{2\sigma^2} \sum (y_i - \mu)^2\right) \times 1 \\ &\propto \exp\left(-\frac{1}{2\frac{\sigma^2}{n}} \sum (y_i - \bar{y})^2\right) \end{aligned}$$

Thus,

$$\mu|y \sim N(\bar{y}, \frac{\sigma^2}{n})$$

This is exactly the same as the estimate from MLE estimate.

NOTE

In case that the variance is unknown,

$$f(\sigma^2) \propto \frac{1}{\sigma^2}$$

This is equivalent to the following:

$$\Gamma^{-1}(0, 1)$$

Thus the posterior for σ^2

$$\sigma^2|y \sim \Gamma^{-1}\left(\frac{n-1}{2}, \frac{1}{2} \sum (y_i - \bar{y})^2\right)$$

11.10 Jeffreys Priors

Jeffreys Prior

$$f(\theta) \propto \sqrt{I(\theta)}$$

For instance,

11.10.1 Gaussian

$$Y_i \sim (\mu, \sigma^2) \rightarrow f(\mu) \propto 1, f(\sigma^2) \propto \frac{1}{\sigma^2}$$

11.10.2 Bernoulli

$$Y_i \sim B(\theta) \rightarrow f(\theta) \propto \theta^{-\frac{1}{2}}(1-\theta)^{-\frac{1}{2}} \sim \text{Beta}(\frac{1}{2}, \frac{1}{2})$$

11.10.3 Side Note: Fisher Information

The Fisher information (for one paramter):

$$I(\theta) = E[(\frac{d}{d\theta} \log(f(X|\theta)))^2]$$

The Expectation is with respect to X with a PDF $f(x|\theta)$.

For instance, suppose the $X|\theta \sim N(\theta, 1)$

$$\begin{aligned} f(x|\theta) &= \frac{1}{\sqrt{2\pi}} \exp[-\frac{1}{2}(x-\theta)^2] \\ \log(f(x|\theta)) &= \log(\frac{1}{\sqrt{2\pi}}) - \frac{1}{2}(x-\theta)^2 \\ \frac{d}{d\theta} \log(f(x|\theta)) &= x - \theta \\ (\frac{d}{d\theta} \log(f(x|\theta)))^2 &= (x - \theta)^2 \\ E[(\frac{d}{d\theta} \log(f(x|\theta)))^2] &= (x - \theta)^2 = \text{Var}(X) = 1 \end{aligned}$$

11.10.4 Prior predictive distribution

Prior distribution vs. prior predictive distribution

<https://stats.stackexchange.com/questions/394648/differences-between-prior-distribution-and-prior-predictive-distribution>

Stochastic Processes

<https://www.youtube.com/watch?v=TuTmC8aOQJE>

11.11 Linear regression

11.11.1 Review

For single covariate of x :

$$E[y] = \beta_0 + \beta_1 x$$

Thus,

$$Y \sim N(\beta_0 + \beta_1 x, \sigma^2)$$

For multiple covariates,

$$E[y] = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_k x_k$$

Side Note It is interesting to think y as the expected value of of the combination of βx .

11.11.2 When σ^2 is known

If we use a Jeffreys prior and assume σ^2 is known, we will get β that has the same mean as the standard OLS. Assuming we only has one covariate:

$$\beta_1 | y \sim N\left(\frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}, \frac{\sigma^2}{\sum_{i=1}^n (x_i - \bar{x})^2}\right)$$

For mutiple covariates, we can use matrix notation,

$$\beta | y \sim N((X^t X)^{-1} X^t y, (X^t X)^{-1} \sigma^2)$$

11.11.3 When σ^2 is unknown

When both β and σ^2 are unknown, the standard prior is the non-informative Jeffreys prior:

$$f(\beta, \sigma^2) \propto \frac{1}{\sigma^2}$$

The posterior mean for β is the same as standard OLS estimates. The posterior for β conditional on σ^2 is the same normal distribution when σ^2 is known. However, the marginal posteior distribution for β , with σ^2 integrated out, is a t distribution.

The t distribution has the mean of $(X^t X)^{-1} X^t y$ and variance matrix, $s^2 (X^t X)^{-1}$, where $s^2 = \sum_{i=1}^n (y_i - \hat{y}_i)^2 / (n - k - 1)$

The variance of σ^2 is an inverse gamma distribution

$$\sigma^2|y \sim \Gamma^{-1}\left(\frac{n-k-1}{2}, \frac{n-k-1}{2}s^2\right)$$

Chapter 12

Bayesian - 2

12.1 Components of Bayesian models

$$y_i = \mu + \epsilon_i$$

Where,

$$\epsilon_i \sim N(0, \sigma^2)$$

$$y_i \sim N(\mu, \sigma^2)$$

(Thus, y_i is = to a fixed μ plus with a ϵ_i , whereas $y_i \sim N(\mu, \sigma^2)$. These two expressions are not exactly the same, but they are connected.)

Likelihood: $P(y|\theta)$ ($P(y, \theta) = P(\theta)P(y|\theta)$)

Prior: $P(\theta)$

Posterior:

$$P(\theta|y) = \frac{P(\theta, y)}{P(y)} = \frac{P(\theta, y)}{\int P(\theta, y)d\theta} = \frac{P(\theta)P(y|\theta)}{\int P(\theta)P(y|\theta)d\theta} = \frac{P(\theta)P(y|\theta)}{\int P(\theta)P(y|\theta)d\theta}$$

Markts

- (1) The only random variables in frequentist models are the data. In contrast, Bayesian paradigm also uses probability to describe one's uncertainty about unknown model parameters.
- (2) Consider the following model for binary outcome y :

$$y_i | \theta_i \sim \text{Bern}(\theta_i), i = 1, 2, 3 \dots 6$$

$$\theta_i | \alpha \sim \text{Beta}(a, b_0), i = 1, 2, 3 \dots 6$$

$$\alpha \sim \text{Exp}(r_0)$$

Thus, the joint distribution of all variable:

$$\prod_{i=1}^6 [\theta_i^{y_i} (1 - \theta_i)^{1-y_i} \frac{\Gamma(a+b_0)}{\Gamma(a)\Gamma(b_0)} \theta_i^{a-1} (1 - \theta_i)^{b_0-1}] r_0 e^{-r_0 \alpha}$$

(Question: Why not write it as a_i ?)

12.2 Monte Carlo Estimation

12.2.1 Sample Mean: Application of Central Limit Theorem

If we simulate 100 samples from a $\text{Gamma}(2,1)$, what is the approximate distribution of the sample average $\bar{x}^* = \frac{1}{100} \sum_{i=1}^{100} x_i^*$?

As we know, based on the central limit theorem, the approximating distribution is normal with the mean equal to the sample mean, and the variance equal to the variance of the original variable divided by the sample size. We know that the mean for $\text{Gamma}(2,1)$ is $\frac{2}{1} = 2$ and variance is $\frac{2}{1^2} = 2$. Thus, we know that the mean for the distribution for \bar{x}^* is 2, whereas the variance is $\frac{2}{100}$. Thus, it is $N(2, 0.02)$.

Thus, we can get the generalized formula as follows:

$$\bar{\theta}^* \sim N(E(\theta), \frac{\text{Var}(\theta)}{m})$$

Note that, it is the variance for $\bar{\theta}^*$, not θ . The approximate for the variance for θ is $\frac{1}{m} \sum (\theta_i^* - \bar{\theta}^*)^2$.

The following R code is for the θ :

```
sample_data_gamma<-rgamma(10000,4,2)
mean(sample_data_gamma)
```

```
## [1] 2.003884
```

```
var(sample_data_gamma)
```

```
## [1] 0.996932
```

The following R code is for the $\bar{\theta}^*$. As we can see, the variance is 0.0016, which is close the theoretical value of $\frac{1}{1000} = 0.001$.


```
set.seed(123)
mean_c<-c()
for(i in 1:10)
{mean_c[i]<-mean(rgamma(1000,4,2))}
var(mean_c)
```

```
## [1] 0.001634891
```

12.2.2 Expected value and probability

If you know that $\theta \sim \text{Beta}(5, 3)$, what is the approximate for the $E(\frac{\theta}{1-\theta})$? You can use the following R code to calculate it.

```
sample_data<-rbeta(1000,5,3)
mean(sample_data/(1-sample_data))
```

```
## [1] 2.789685
```

If you want to calculate the approximate for the probability that $\frac{\theta}{1-\theta}$ is greater 1.

```
mean((sample_data/(1-sample_data))>1)
```

```
## [1] 0.789
```

12.2.3 Quantile

Use Monte Carlo to approximate the 0.3 quantile of $N(0,1)$. Note that, the idea of quantile is to quantify the probability. Thus, the number we will get for 0.3 quantile if the value for the random variable's cdf $\int_{-\infty}^{\text{quantile-number}} dx$. As we can see below, `quantile(sample_data2,0.3)` gets the result of -0.52, which is the same from the `qnorm(0.3)`. Note that, the cdf of `pnorm(-0.52)` will get the probability of 0.3.

```
sample_data2<-rnorm(10000,0,1)
quantile(sample_data2,0.3)
```

```
##          30%
## -0.5260847
```

```
qnorm(0.3)
```

```
## [1] -0.5244005
```

```
pnorm(-0.52)
```

```
## [1] 0.3015318
```

12.3 Prior predictive distributions

$$y|\theta \sim \text{Bin}(10, \theta)$$

$$\theta \sim \text{Beta}(2, 2)$$

Simulate: (1) θ^* from beta (2) Given θ , draw $y_i^* \sim \text{Bin}(10, \theta_i^*)$ (3) Get pairs (y_i^*, θ_i^*)

Chapter 13

Trying

13.1 The Basic Idea

$$L(\beta, D|Y) = \int \prod_{i=1}^n f_{y_i|u}(y_i|b, \beta) f_{b_i}(b_i|D) db_i$$

Notations :

y : Variable for the fixed effect

b : Variable for the random effect

β : Parameters for the fixed effect

D : Parameters for the random effect

The dimension of the integral is equal to the levels of the random factors (i.e., the number of observations).

13.2 Model and R Code

Covariance Matrix for n observations:

$$V = \sigma^2 \begin{bmatrix} 1 & \rho & \rho^2 & \dots & \rho^{n-1} & \rho^n \\ \rho & 1 & \rho & \dots & \rho^{n-2} & \rho^{n-1} \\ \rho^2 & \rho & 1 & \dots & \rho^{n-3} & \rho^{n-2} \\ \dots & & & & & \\ \rho^n & \rho^{n-1} & \rho^{n-2} & \dots & \rho & 1 \end{bmatrix}$$

The inverse matrix is as follows:

$$Q = V^{-1} = \frac{1}{\sigma^2(1-\rho)} \begin{bmatrix} 1 & -\rho & 0 & \dots & 0 & 0 \\ -\rho & 1+\rho^2 & -\rho & \dots & 0 & 0 \\ 0 & -\rho & 1+\rho^2 & \dots & 0 & 0 \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 0 & 0 & 0 & \dots & 1+\rho^2 & -\rho \\ 0 & 0 & 0 & \dots & -\rho & 1 \end{bmatrix}$$

$$N(-\sum_{j \neq k} Q_{kj} b_j^{(m)} Q_{kk}^{-1}, Q_{kk}^{-1})$$

$$\ln L(\beta, \theta; Y, b) = \ell = \ln f_{Y|b}(Y|b, \beta) + \ln f_b(b|\theta)$$

$$a^{(m+1)} = a^{(m)} + \tau(a^{(m)})^{-1} S(a^{(m)})$$

Where,

$$\tau(a) = -E(\frac{\partial^2 \ell}{\partial \alpha \partial \alpha'} | Y)$$

$$S(a) = E(\frac{\partial \ell}{\partial \alpha} | Y)$$

Note that, α is a combination of two sets of parameters.

$$\alpha = \begin{pmatrix} \beta \\ b \end{pmatrix}$$

$$\ell = \sum_{i=1}^n \{ [y_i \ln(\frac{e^{\beta^T x_i + b_i}}{1 + e^{\beta^T x_i + b_i}}) + (1 - y_i) \ln(1 - \frac{e^{\beta^T x_i + b_i}}{1 + e^{\beta^T x_i + b_i}})] + \ln f_b(b_i|\theta) \}$$

$$\frac{\partial \ln f(Y|b, \beta)}{\partial \beta} = X'(Y - E(Y|b))$$

$$\frac{\partial \ln f(Y|b, \beta)}{\partial \beta \partial \beta'} = -X'(Y - E(Y|b))$$

$$\begin{aligned} \nabla \ell &= \sum_{i=1}^n [y_i \frac{1}{p(\beta^T x_i + b_i)} \frac{\partial p(\beta^T x_i + b_i)}{\partial (\beta^T x_i + b_i)} \frac{\partial (\beta^T x_i + b_i)}{\partial \beta} + (1 - y_i) \frac{1}{1 - p(\beta^T x_i + b_i)} (-1) \frac{\partial p(\beta^T x_i + b_i)}{\partial (\beta^T x_i + b_i)} \frac{\partial (\beta^T x_i + b_i)}{\partial \beta}] \\ &= \sum_{i=1}^n x_i^T [y_i - p(\beta^T x_i + b_i)] \\ &= \sum_{i=1}^n x_i^T [y_i - \frac{e^{\beta^T x_i + b_i}}{1 + e^{\beta^T x_i + b_i}}] \end{aligned}$$

The Newton Raphson algorithm needs the second order.

$$\begin{aligned}
 \nabla^2 \ell &= \frac{\partial \sum_{i=1}^n x_i^T [y_i - p(\beta^T x_i + b_i)]}{\partial \beta} \\
 &= - \sum_{i=1}^n x_i^T \frac{\partial p(\beta^T x_i + b_i)}{\partial \beta} \\
 &= - \sum_{i=1}^n x_i^T \frac{\partial p(\beta^T x_i + b_i)}{\partial (\beta^T x_i + b_i)} \frac{\partial (\beta^T x_i + b_i)}{\partial \beta} \\
 &= - \sum_{i=1}^n x_i^T p(\beta^T x_i + b_i) (1 - p(\beta^T x_i + b_i)) x_i \\
 &= - \sum_{i=1}^n x_i^T \frac{e^{\beta^T x_i + b_i}}{1 + e^{\beta^T x_i + b_i}} \left(1 - \frac{e^{\beta^T x_i + b_i}}{1 + e^{\beta^T x_i + b_i}}\right) x_i
 \end{aligned}$$

Using the Newton Raphson, the following code calculates the basic logistic model, without any random effects. As we can see, it produces the same result as the R generic function of GLM. Note that, the function of `Expit` and the variables of `y` and `x` are from the last block of R code.

```

#n<-10
x_intercept<-rep(1,n)
x_intercept<-as.matrix(x_intercept)

tolerance=1e-3
max_its=2000;iteration=1;difference=2
W<-matrix(0,n,n)
beta_old<-0.4

while(difference>tolerance & iteration<max_its)
{
  # The first order
  f_firstorder<-t(x_intercept)%*(y-Expit(x_intercept%*beta_old))

  # The second order
  diag(W) = Expit(x_intercept%*beta_old)*(1-Expit(x_intercept%*beta_old))

  f_secondorder<--t(x_intercept)%*W%*x_intercept

  # Calculate the beta_updated
  beta_updated=beta_old-(solve(f_secondorder)%*f_firstorder)

  difference=max(abs(beta_updated-beta_old));

  iteration=iteration+1;

```

```

    beta_old=beta_updated}

beta_old

glm(y~1, family=binomial)$coefficients

#install.packages("CVTuningCov")
library(CVTuningCov) # Will be used to generate AR1 matrix

set.seed(123)
y<-c(1,1,1,0,0,1,1,0,1,0) ## observations

n=length(y) # the number of observations

#Establish the exp function
Expit<-function(x){exp(x)/(1+exp(x))}
#Y: observations
#b: random effect
#beta_0:fixed effect->intercept (or, mean of Y)

log_pdf_function<-function(Y,b,beta_0)
{mean_prob<-Expit(beta_0+b)
  dbinom(Y,1,mean_prob,log = TRUE)
}

b_records<-rep(0,n) #Initial values for the random effect
rho_records<-0.5 #Initial value for rho
sigma_records<-2 #Initial value for sigma
mean_0<-0 # Initial mean value for normal distribution (of the random effect)
beta<-0.5 # Initial value for the intercept of Y

f_random<-function(sigma_records, rho_records,beta)
{

co_matrix<-(sigma_records^2)*AR1(n,rho_records) # covariance matrix
co_matrix_inverse<-solve(co_matrix) # inverse covariance matrix

for (k in 1:n)
{
  # Variance for the random effect
  sd_0<-1/(co_matrix_inverse[k,k])

```

```

for(j in 1:n)
{ # Make sure that k is not equal to j, otherwise 0
  Q_kj<-ifelse(j!=k,co_matrix_inverse[k,j],0)
  # Calculate the mean for the random effect; sum of mean in a loop
  mean_0<-mean_0-(Q_kj/co_matrix_inverse[k,k])*b_records[j]
}
# Draw a random number from the normal distribution for the random effect
b_candidate<-rnorm(1,mean_0,sd_0)

current_lp<-log_pdf_function(y[k],b_records[k],beta)
candidate_lp<-log_pdf_function(y[k],b_candidate,beta)

Smaller_value<-min(exp(candidate_lp-current_lp),1)
# Draw a random number from the uniform distribution
Random_probability<-runif(1)
# Update b (i.e., random variable)
b_records[k]<-ifelse(Random_probability<Smaller_value,b_candidate,b_records[k])
}

return(b_records)
}

# Print result
#b_records<-f_random(1,0.8,0.3)

```

In the following, I will try to add the random effect.

```

x_intercept<-rep(1,n)
x_intercept<-as.matrix(x_intercept)

#We need to set random starting values.

tolerance=1e-3
max_its=2000;iteration=1;difference=2
W<-matrix(0,n,n)
beta_old<-0.4

while(difference>tolerance & iteration<max_its)
{
  b_records<-f_random(2,0.9,beta_old)
  print("first")
  print(beta_old)
  print(b_records)
  # The first order
  f_firstorder<-t(x_intercept)%*(y-Expit(x_intercept%*beta_old+b_records))
  print(f_firstorder)
}

```

```

# The second order
diag(W) = Expit(x_intercept**beta_old+b_records)*(1-Expit(x_intercept**beta_old+b_

f_secondorder<--t(x_intercept)**W**x_intercept

# Calculate the beta_updated
beta_updated=beta_old-(solve(f_secondorder)**f_firstorder)

difference=max(abs(beta_updated-beta_old));

iteration=iteration+1;

beta_old=beta_updated
}

beta_old

```

13.3 glmmTMB package

```

# https://becarioprecario.bitbucket.io/inla-gitbook/ch-intro.html
#https://cran.r-project.org/web/packages/glmmTMB/vignettes/covstruct.html
#install.packages("glmmTMB")
library(glmmTMB)

times <- factor(1:n)
levels(times)
group <- factor(rep(1,n))
dat0 <- data.frame(y,times,group)

glmmTMB(y ~ ar1(times + 0 | group), data=dat0)

```